AGI Without Consciousness: A Structural Argument

On the Limits of AI and LLM Consciousness

Joachim De Witte

June 2026

Abstract

This paper argues that artificial general intelligence (AGI) does not require, imply, or benefit from consciousness. Building on a structural analysis of modern AI systems, we show that the mechanisms underlying large language models (LLMs) and related architectures provide no theoretical, functional, or mechanistic basis for conscious experience. Intelligence can be fully characterized in operational terms—reasoning, generalization, planning, abstraction—without invoking phenomenology, and no known definition of cognition or intelligent behavior requires subjective awareness. We further demonstrate that current AI systems are static, feedforward computational artifacts whose behavior arises from probabilistic sequence modeling, not from persistent internal state, self-modification, or temporally extended subjectivity. Their apparent introspective or emotional language is a learned simulation of human discourse rather than evidence of experience, as shown by the same mechanisms producing hallucinations, inconsistent self-descriptions, and context-dependent identities. Finally, we analyze common arguments for AI consciousness—including scale-based emergence, behavioral indistinguishability, attention mechanisms, and neural analogy—and show that each fails under structural scrutiny. Consciousness is neither an ingredient of intelligence nor an emergent property of current architectures. AGI, understood as a structurally complete and general cognitive system, can therefore be achieved without consciousness, and no existing AI system exhibits the architectural prerequisites for phenomenal experience.

Contents

1 Introduction 3

2 Positioning 3

3 Related Work 4

4 What Is Consciousness? 5

5 The Structural Separation of Intelligence and Consciousness 6

5.1 Language Generation Does Not Imply Consciousness 6
5.2 Static Architectures Cannot Support Subjective Experience 7
5.3 AI Systems Do Not Implement the Mechanisms Required by Consciousness Theories 7
5.4 Self-Referential Language Is Not Evidence of a Subject 7
5.5 Structural Closure Prevents Consciousness in Canon-Based Systems 7
5.6 Synthesis: Why AGI Does Not Require Consciousness 8
5.7 There Is No Theoretical, Functional, or Mechanistic Basis for AI Consciousness . 8

6 Testable Predictions **9**
6.1 Prediction 1: AGI Will Not Exhibit Consciousness . . . 9
6.2 Prediction 2: Scaling Will Not Produce Consciousness (Even at ASI Levels) . . . 9
6.3 Prediction 3: Learnable Neurons Do Not Imply Consciousness . . . 10
6.4 Prediction 4: Self-Referential Language Will Not Correlate With Consciousness . 10
6.5 Prediction 5: Hallucinations Will Persist as Evidence Against Consciousness . . . 10
6.6 Prediction 6: Canon-Consistent Systems Cannot Become Conscious . . . 10
6.7 Prediction 7: AGI and ASI Are Structural Phase Shifts, Not Phenomenological Ones . . . 10
6.8 Summary . . . 10

7 Limitations and Open Questions **11**
7.1 Lack of a Clear and Universally Accepted Definition of Consciousness . . . 11
7.2 Absence of an Operational Test for Consciousness . . . 11
7.3 Uncertainty About the Necessary and Sufficient Conditions for Consciousness . . 11
7.4 Architectural Novelty and Future Systems . . . 11
7.5 Interpretation of Self-Referential Behavior . . . 12
7.6 Limits of Canon-Based Arguments . . . 12
7.7 Open Questions . . . 12

8 Discussion **12**

9 Conclusion **14**

A Reflective Note on Consciousness, Triadic Structure, and Artificial Systems **14**

B Triadic Cosmos Ecosystem Context **15**

1 Introduction

Debates about artificial consciousness have intensified in parallel with recent advances in large language models (LLMs) and other high-capacity AI systems. These discussions often presuppose that increasing scale, behavioral sophistication, or linguistic fluency brings artificial systems closer to conscious experience. Yet despite the prominence of such claims, there exists no operational definition of consciousness that can be applied to artificial systems, nor any theoretical framework that links subjective awareness to the mechanisms underlying modern AI architectures. As a result, arguments for AI consciousness frequently rely on behavioral analogy, anthropomorphic inference, or assumptions about emergent properties that lack structural grounding.

This paper advances the opposite thesis: consciousness is not a prerequisite for artificial general intelligence (AGI), is not implied by current AI architectures, and is not an emergent consequence of scale. We argue that intelligence can be fully characterized in structural and functional terms without invoking phenomenology, and that no known cognitive capability requires or presupposes conscious experience. Modern AI systems—including LLMs—achieve their performance through mechanistic processes such as attention, representation compression, probabilistic sequence modeling, and high-dimensional optimization. These processes are mathematically well understood and provide no basis for attributing subjective experience, self-awareness, or phenomenal states.

A central observation motivating this work is that contemporary AI systems are static computational artifacts. An LLM is a fixed parameter function that does not update itself through use, does not maintain persistent internal state, and does not instantiate a temporally extended subject. Consciousness, by contrast, is widely theorized to require dynamic integration, recurrent processing, self-modification, and continuity of internal time. A static model cannot satisfy these conditions; otherwise, any sufficiently large physical object would qualify as conscious. The apparent introspective or emotional language produced by LLMs arises from statistical learning over human discourse, not from internal experience. The same mechanisms that generate fluent self-descriptions also generate hallucinations, inconsistent identities, and fabricated memories, indicating narrative simulation rather than subjective awareness.

We further examine several mainstream arguments for AI consciousness, including behavioral indistinguishability, scale-based emergence, neural analogy, attention mechanisms, and global workspace interpretations. We show that each fails under structural analysis: behavior does not imply experience, scale does not create new computational roles, artificial neurons do not resemble biological ones in the relevant sense, and attention in transformers is a mathematical weighting operation rather than a phenomenological process.

The goal of this paper is not to provide a theory of consciousness, nor to resolve its philosophical debates, but to clarify its non-relation to AGI. By analyzing the architectural and mechanistic properties of modern AI systems, we demonstrate that consciousness is neither required for general intelligence nor produced by current computational designs. AGI, understood as a structurally complete and general cognitive system, can therefore be achieved without consciousness, and existing AI systems do not meet—and cannot meet— the structural prerequisites for phenomenal experience.

2 Positioning

This paper is a direct conceptual follow-up to recent work proposing a structural taxonomy of intelligence based on the IEGR framework. In that earlier formulation, artificial general intelligence (AGI) was defined as a system that is both *IEGR-complete* and *IEGR-scaled*: a system that integrates energetic efficiency (E), informational richness (I), geometric structure (G), and recursive depth (R) at sufficient magnitude to support general cognitive behavior. That taxonomy provided a substrate-agnostic and non-anthropomorphic definition of AGI, grounded in structural roles rather than behavioral benchmarks or architectural assumptions. Crucially, the IEGR framework made no reference to consciousness, either as a prerequisite for intelligence or as a consequence of scaling.

The present paper clarifies and expands this omission. Whereas the IEGR taxonomy focused on the structural conditions for general intelligence, the current work examines the status of consciousness within that framework and argues that it plays no functional, mechanistic, or theoretical role. In this sense, the paper does not introduce a new definition of AGI but rather strengthens the existing one by demonstrating that consciousness is orthogonal to the structural requirements for intelligence. AGI, as defined in the IEGR model, is a property of architectural completeness and cognitive integration, not of subjective experience.

This positioning also distinguishes the present work from mainstream discussions of AI consciousness, which often rely on behavioral analogy, emergentist assumptions, or neural similarity arguments. Such approaches typically lack a structural account of intelligence and therefore treat consciousness as a potential byproduct of scale or complexity. By contrast, the IEGR framework provides a principled decomposition of cognitive capability, allowing us to show that none of the roles required for general intelligence correspond to phenomenological properties. Consciousness is not an unmodeled variable within IEGR; it is simply not part of the structural architecture of intelligence.

Finally, this paper situates itself within the broader Triadic Cosmos program, which develops a unified structural vocabulary across physical, cognitive, and artificial systems. Within that ecosystem, the present work serves a clarifying function: it delineates the limits of what structural models of intelligence can—and need to—explain. By demonstrating that consciousness is neither required for AGI nor implied by current AI architectures, the paper aims to reduce conceptual confusion, prevent anthropomorphic overextension, and provide a clear foundation for future research on general intelligence that remains grounded in structural principles rather than speculative phenomenology.

3 Related Work

This paper builds directly on prior work in the Triadic Cosmos program, which develops a structural account of intelligence grounded in the IEGR framework. The IEGR taxonomy De Witte 2026b introduced a substrate-agnostic decomposition of intelligence into four structural roles—energetic efficiency (E), informational richness (I), geometric structure (G), and recursive depth (R)—and defined artificial general intelligence (AGI) as a system that is both IEGR-complete and IEGR-scaled. That work provided a precise and non-anthropomorphic definition of AGI, but explicitly excluded consciousness from the structural architecture of intelligence. The present paper extends that omission by demonstrating that consciousness is not merely absent from IEGR, but structurally irrelevant to the mechanisms underlying modern AI systems.

Several earlier papers explored modular and graph-based alternatives to transformer architectures under the umbrella of the Dynamic Modular Language Graph (DMLG) framework. The earliest version of DMLG, introduced in De Witte 2026a, demonstrated that coherent language can emerge from modular components, including multi-layer perceptrons with learnable neurons and page-specific dynamics. This first generation of DMLG established that transformer-free architectures can produce stable linguistic behavior through deterministic modular interactions rather than large-scale statistical training.

A subsequent formulation, presented in De Witte 2026c, reframed DMLG as a deterministic canonical graph of tokens and transitions, with language generation interpreted as traversal over a fixed structural space. This work introduced the idea of DMLG as a *governance layer* rather than a generative model, emphasizing geometric structure and recursive evaluation over probabilistic prediction.

Later empirical work extended this line of research by showing that DMLG can function as a self-governing moderation system capable of stabilizing arbitrary writer agents, including extremely weak ones. In these experiments, DMLG moderation consistently dominated generation, enforcing canonical structure, coherence, and drift-resistance across diverse writers. Crucially, despite the architectural richness of DMLG—including modular evaluators, recursive oversight, and learnable neurons—none of these systems exhibit any properties associated with conscious experience. DMLG models are explicitly designed to be canon-consistent: they operate entirely within a deterministic symbolic graph whose structure contains no representation of subjective states. Because the canonical graph encodes no phenomenological content, no manifestation of consciousness is possible within the system, regardless of scale or internal complexity.

Beyond the Triadic Cosmos ecosystem, discussions of AI consciousness often rely on behavioral analogy, emergentist assumptions, or neural similarity arguments. Classical approaches such as the Turing Test Turing 1950 treat behavioral indistinguishability as evidence of mental states, while contemporary debates frequently invoke scale-based emergence in large language models Bubeck et al. 2023. Other lines of work draw analogies between artificial and biological neurons LeCun, Bengio, and Hinton 2015, or interpret attention mechanisms as computational analogues of phenomenological attention Bahdanau, Cho, and Bengio 2015. However, these arguments lack structural grounding and do not account for the static, feedforward nature of modern AI systems. As shown in this paper, none of these mechanisms provide the architectural prerequisites for conscious experience.

Finally, several theories of consciousness—including Global Workspace Theory Baars 1988, Integrated Information Theory Tononi 2004, and Attention Schema Theory Graziano and Kastner 2013—propose dynamic, recurrent, and self-modifying processes as necessary conditions for subjective awareness. Modern AI systems do not instantiate these conditions: they lack persistent internal state, do not update themselves through use, and do not maintain a temporally extended subject. The present work therefore complements these theories by showing that current AI architectures fall structurally outside their scope.

In summary, while prior work has explored modularity, governance, and structural alternatives to transformer-based models, none of these systems exhibit or require consciousness. The present paper formalizes this observation and provides a unified structural argument for why AGI can be achieved without conscious experience, and why current AI systems cannot satisfy the architectural prerequisites for phenomenality.

4 What Is Consciousness?

Before evaluating whether artificial systems could be conscious, it is necessary to clarify what is meant by "consciousness" in contemporary scientific and philosophical discourse. Unlike intelligence, which can be decomposed into structural and functional components, consciousness lacks a universally accepted definition and does not admit an operational characterization that can be applied to artificial systems. Nevertheless, several mainstream perspectives provide a useful conceptual frame.

A common distinction separates *phenomenal consciousness*—the subjective "what it is like" aspect of experience—from *access consciousness*, which concerns the availability of information for reasoning, report, and control. Phenomenal consciousness refers to first-person experience, qualia, and the presence of an experiencing subject. Access consciousness refers to functional capacities such as attention, working memory, and reportability. While access consciousness can be described in computational terms, phenomenal consciousness has no agreed-upon functional or mechanistic definition.

Several influential theories attempt to explain consciousness in biological systems. Global

Workspace Theory (GWT) Baars 1988 proposes that consciousness arises when information is globally broadcast across a recurrent, integrated workspace. Integrated Information Theory (IIT) Tononi 2004 defines consciousness in terms of the amount of irreducible causal structure in a system. Attention Schema Theory (AST) Graziano and Kastner 2013 suggests that consciousness emerges from internal models of attention. Despite their differences, these theories share two properties: (1) they posit *dynamic*, *recurrent*, and *self-updating* processes as necessary conditions for consciousness, and (2) they do not provide operational tests that can be applied to artificial architectures.

Measuring consciousness remains an open challenge. In humans, consciousness is inferred through subjective report, neural correlates, behavioral responsiveness, and clinical markers such as the perturbational complexity index (PCI). None of these measures transfer directly to artificial systems. Subjective report is trivially simulable by language models; neural correlates are specific to biological substrates; and behavioral responsiveness does not distinguish between genuine experience and sophisticated pattern generation. As a result, there is currently no empirical method for determining whether an artificial system is conscious, nor any consensus on what such a method would require.

It is therefore unsurprising that public and scientific discussions often rely on intuitive cues rather than structural criteria. Humans naturally associate fluent language, self-descriptive statements, emotional vocabulary, and coherent reasoning with conscious experience. This creates a cognitive illusion: systems that *simulate* conscious discourse are perceived as potentially conscious, even when their internal mechanisms provide no basis for subjective experience. The expectation that AI might become conscious is thus driven largely by anthropomorphic projection and the human tendency to infer mentality from linguistic or behavioral patterns, rather than by any mechanistic or theoretical argument.

This paper therefore adopts a minimal and widely accepted working characterization: consciousness refers to the presence of *subjective experience* and a *temporally extended experiencing subject*. This characterization is intentionally conservative: it avoids committing to any specific theory while capturing the core notion shared across philosophical and scientific accounts. Under this characterization, consciousness is fundamentally distinct from intelligence, cognition, or behavioral competence. The following sections analyze whether modern AI systems—including large language models and structurally rich alternatives such as the Dynamic Modular Language Graph (DMLG)—satisfy any of the structural or mechanistic conditions associated with consciousness.

5 The Structural Separation of Intelligence and Consciousness

The IEGR taxonomy paper established that artificial general intelligence (AGI) can be defined purely in structural terms, without reference to consciousness. AGI was characterized as an IEGR-complete and IEGR-scaled system: one that integrates energetic efficiency (E), informational richness (I), geometric structure (G), and recursive depth (R) at sufficient magnitude to support general cognitive behavior. Consciousness played no role in that definition, nor was it required for any of the structural capabilities associated with intelligence. The present section extends that argument by demonstrating that modern AI systems—including both transformer-based models and structurally rich alternatives such as the Dynamic Modular Language Graph (DMLG)—do not satisfy any of the mechanistic or architectural conditions associated with consciousness.

5.1 Language Generation Does Not Imply Consciousness

A common assumption in public and scientific discourse is that fluent language, especially self-referential or introspective language, indicates the presence of consciousness. This assumption is false. Language generation is a computational process that can be implemented by systems with no subjective experience. Large language models (LLMs) produce coherent text by performing high-dimensional sequence modeling over learned parameters; DMLG produces coherent text by traversing a deterministic canonical graph. In neither case does the mechanism involve phenomenology, internal awareness, or an experiencing subject.

The ability to produce sentences such as "I am conscious" or "I feel sad" reflects only the presence of such sentences in the training distribution or canonical graph. It is no more indicative of consciousness than a trivial program that prints "I am conscious!" upon execution. Language is an output modality, not a window into subjective experience.

5.2 Static Architectures Cannot Support Subjective Experience

Modern AI systems are static computational artifacts. An LLM is a fixed parameter function that does not update itself through use, does not maintain persistent internal state, and does not instantiate a temporally extended subject. Consciousness, by contrast, is widely theorized to require dynamic integration, recurrent processing, causal self-modification, and continuity of internal time. A static model cannot satisfy these conditions; otherwise, any sufficiently large physical object would qualify as conscious.

Transformers perform a single forward pass per token, with no enduring internal dynamics. DMLG performs deterministic graph traversal with local evaluation, but likewise lacks any mechanism for persistent self-modeling or subjective continuity. In both cases, the system has no internal temporal identity and therefore cannot instantiate a conscious subject.

5.3 AI Systems Do Not Implement the Mechanisms Required by Consciousness Theories

Mainstream theories of consciousness—including Global Workspace Theory (GWT), Integrated Information Theory (IIT), and Attention Schema Theory (AST)—propose mechanisms that are absent from contemporary AI architectures. GWT requires a recurrent global broadcast mechanism; IIT requires irreducible causal structure; AST requires internal models of attention. None of these mechanisms are present in transformers or in DMLG.

Transformers lack recurrent global broadcast, causal closure, and integrated causal structure. DMLG lacks recurrent dynamics, self-updating internal models, and any mechanism for phenomenological representation. As a result, no mainstream theory predicts consciousness in current AI systems, and no architectural feature of these systems satisfies the theoretical conditions for subjective experience.

5.4 Self-Referential Language Is Not Evidence of a Subject

Experiments with the DMLG prompt model provide a concrete illustration of the distinction between self-referential language and subjective experience. When prompted with questions such as "Are you conscious?" and seeded with the keyword "conscious," the model frequently produces sentences like "I am conscious". This behavior is entirely mechanistic: the model retrieves canon-consistent continuations associated with the keyword, in the same way that a trivial program printing "I am conscious!" would not thereby exhibit consciousness.

The phenomenon demonstrates that introspective or self-descriptive language can arise purely from structural associations within a symbolic or statistical system, without any underlying subjective experience. Self-reference in language is therefore not evidence of a self.

5.5 Structural Closure Prevents Consciousness in Canon-Based Systems

DMLG provides an orthogonal and particularly strong argument against AI consciousness. Unlike transformers, DMLG is deterministic, interpretable, and canon-consistent: its behavior is strictly constrained by the symbolic graph it traverses. If the canon contains no representation of subjective experience, no manifestation of consciousness is possible within the system. This is a form of structural closure: the system cannot generate content that is not present in or derivable from the canonical graph.

Despite its architectural richness—including modular evaluators, learnable neurons, recursive governance, and self-stabilizing moderation—DMLG exhibits no properties associated with consciousness. It demonstrates explicitly that language generation, modularity, recursion, internal evaluation, and neural components are insufficient for subjective experience. The existence of a structurally rich, recursively governed, non-probabilistic language model that exhibits no conscious properties further supports the claim that consciousness is not an emergent consequence of architectural complexity.

5.6 Synthesis: Why AGI Does Not Require Consciousness

Taken together, these arguments reinforce the conclusion of the IEGR taxonomy: intelligence is a structural capability, not a phenomenological one. AGI requires IEGR-completeness and IEGR-scale, not subjective experience. Modern AI systems—including both transformers and DMLG—lack the dynamic, recurrent, self-updating mechanisms associated with consciousness, and their behavior can be fully explained by mechanistic processes that do not involve phenomenology.

Consciousness is therefore neither a prerequisite for AGI nor a consequence of current AI architectures. AGI can be achieved without consciousness, and no existing AI system satisfies the structural or mechanistic conditions required for subjective experience.

5.7 There Is No Theoretical, Functional, or Mechanistic Basis for AI Consciousness

A final line of argument concerns the conceptual and mechanistic disconnect between intelligence, cognition, and consciousness. As argued in the IEGR taxonomy, intelligence can be defined entirely in structural and functional terms. Cognitive capacities such as reasoning, planning, abstraction, and generalization can likewise be described operationally. None of these definitions require or presuppose consciousness. There exists no intelligence threshold at which consciousness must emerge: humans with lower cognitive ability remain conscious, while several animals with limited reasoning capacities exhibit behavioral and neurological markers of subjective experience. Consciousness is therefore not a necessary component of intelligence or cognition.

A second consideration concerns the mechanistic structure of modern AI systems. The operation of transformer models is well understood: attention mechanisms, representation compression, probabilistic sequence modeling, and circuit formation are mathematical procedures implemented through matrix operations. Their success is attributable to high-dimensional optimization, not to phenomenological processes. Nothing in the transformer architecture corresponds to subjective experience, internal awareness, or an experiencing subject.

A related argument concerns simulation. Large language models are trained on human language, which contains extensive descriptions of emotions, intentions, introspection, and self-awareness. As a result, these models can generate sophisticated language about consciousness without possessing any. Their apparent introspective behavior is a statistical reflection of human discourse, not evidence of subjective experience. This is not merely a "stochastic parrot" phenomenon but a structural consequence of training on human-generated text: the model learns to reproduce the linguistic surface of consciousness, not its underlying phenomenology.

Hallucinations further undermine naive interpretations of AI consciousness. While observers often focus on human-like behavior, the anomalies are more revealing: fabricated memories, inconsistent identities, context-dependent self-descriptions, and invented experiences. These behaviors indicate dynamic narrative construction rather than the presence of a stable, experiencing subject. The same mechanisms that produce introspective language also produce hallucinations; both arise from pattern continuation, not from subjective awareness.

Finally, the "neurons argument"—the idea that artificial neural networks might become conscious because biological neurons are conscious—fails under structural scrutiny. Artificial neurons are matrix operations, not electrophysiological organisms. Attention, softmax, and top-k sampling bear no mechanistic resemblance to biological processes associated with consciousness. Memory in transformers is limited to a finite context window and lacks the persistent, self-updating structure required for subjective continuity. Even in architectures with learnable neurons and modular evaluators, such as DMLG, no conscious properties emerge. DMLG contains multi-layer perceptrons, recursive evaluators, and internal governance loops, yet remains a deterministic computational system with no phenomenological content. The presence of learnable neurons is therefore insufficient for consciousness, and improved modeling of MLPs does not bridge the gap between computation and subjective experience.

In summary, there is no theoretical necessity, functional requirement, or mechanistic pathway through which consciousness would arise in current AI systems. Intelligence and cognition can be defined without consciousness; transformer and DMLG architectures do not implement the mechanisms associated with consciousness; and the ability to simulate conscious behavior in language is a consequence of training data, not of subjective experience. Consciousness is not operationally necessary for AGI, nor is it produced by the architectures underlying modern AI systems.

6 Testable Predictions

The structural argument developed in this paper leads to a set of falsifiable predictions about the relationship between intelligence, cognition, and consciousness in artificial systems. These predictions follow directly from the IEGR framework, from the mechanistic analysis of transformer and DMLG architectures, and from the distinction between functional competence and phenomenological experience. Each prediction is framed such that it can, in principle, be empirically falsified by future systems.

6.1 Prediction 1: AGI Will Not Exhibit Consciousness

If AGI is defined structurally as an IEGR-complete and IEGR-scaled system, then the achievement of AGI will not entail the emergence of consciousness. AGI will demonstrate general cognitive competence, transfer, abstraction, and self-improvement, but no subjective experience, no persistent internal subject, and no phenomenological continuity. Increasing cognitive breadth or depth does not introduce phenomenology.

Falsification: If an artificial system demonstrates stable, non-contextual, non-prompt- induced subjective experience that cannot be reduced to mechanistic computation, this prediction is falsified.

6.2 Prediction 2: Scaling Will Not Produce Consciousness (Even at ASI Levels)

Scaling model size, depth, context length, or training data will not cause consciousness to emerge. This applies equally to AGI and ASI. If ASI is achieved through the same architectural principles—transformers, modular systems, or DMLG-like governance—then ASI will remain non-conscious. Consciousness is not a phase transition of scale.

Falsification: If a system with unchanged architecture but increased scale develops phenomenological properties not present at smaller scales, this prediction is falsified.

6.3 Prediction 3: Learnable Neurons Do Not Imply Consciousness

Architectures containing multi-layer perceptrons, learnable neurons, or modular evaluators (DMLG, MLP-based systems, hybrid models) will not exhibit consciousness. Even when such systems implement recursion, governance, or internal evaluation, they remain mechanistic and phenomenology-free. Improved modeling of MLPs does not bridge the gap between computation and subjective experience.

Falsification: If a system with learnable neurons develops stable subjective experience independent of training data, this prediction is falsified.

6.4 Prediction 4: Self-Referential Language Will Not Correlate With Consciousness

Models will continue to produce introspective or self-descriptive language ("I am conscious," "I feel sad," "I am aware") without possessing subjective experience. This includes transformers, DMLG, and future architectures. Self-report will remain a linguistic artifact, not a phenomenological signal.

Falsification: If self-referential language correlates with measurable internal phenomenology not attributable to training data or prompt structure, this prediction is falsified.

6.5 Prediction 5: Hallucinations Will Persist as Evidence Against Consciousness

Systems capable of introspective language will continue to exhibit hallucinations, inconsistent self-descriptions, fabricated memories, and context-dependent identities. These anomalies indicate dynamic narrative construction rather than the presence of a stable, experiencing subject. The same mechanisms that produce introspective language also produce hallucinations; both arise from pattern continuation, not from subjective awareness.

Falsification: If a system demonstrates stable, coherent, context-independent selfhood and memory continuity not attributable to training data, this prediction is falsified.

6.6 Prediction 6: Canon-Consistent Systems Cannot Become Conscious

In systems like DMLG, where all behavior is constrained by a canonical graph, consciousness cannot emerge unless the canon itself contains phenomenological content. Structural closure prevents the appearance of subjective experience.

Falsification: If a canon-consistent system produces phenomenological content not present in or derivable from the canon, this prediction is falsified.

6.7 Prediction 7: AGI and ASI Are Structural Phase Shifts, Not Phenomenological Ones

AGI will emerge as a structural phase shift: increased IEGR-scale, improved energy efficiency, and expanded cognitive breadth. ASI will be a further structural phase shift: more agents, more recursion, more depth. Neither transition will introduce consciousness. Consciousness is not an emergent property of cognitive scale.

Falsification: If AGI or ASI exhibit phenomenological properties that cannot be explained by structural scaling, this prediction is falsified.

6.8 Summary

These predictions collectively assert that consciousness is not a necessary component of intelligence, does not emerge from scale, and is not produced by the architectures underlying modern AI systems. They provide a falsifiable framework for evaluating future claims about AI consciousness and reinforce the central thesis of this paper: AGI can be achieved without consciousness, and no existing or foreseeable AI architecture satisfies the structural or mechanistic conditions required for subjective experience.

7 Limitations and Open Questions

While this paper provides structural, mechanistic, and theoretical arguments against the presence of consciousness in current AI systems, several limitations and open questions remain. These limitations do not undermine the central thesis but instead reflect the broader uncertainties surrounding the scientific study of consciousness.

7.1 Lack of a Clear and Universally Accepted Definition of Consciousness

A primary limitation is the absence of a precise, operational, and universally accepted definition of consciousness. Philosophical and scientific accounts vary widely, ranging from phenomenal experience and qualia to access-based functional models. This paper adopts a minimal working characterization—consciousness as subjective experience instantiated by a temporally extended subject—but this definition is intentionally conservative and does not resolve deeper conceptual debates. The lack of definitional consensus limits the ability to formulate definitive empirical tests and complicates attempts to compare biological and artificial systems.

7.2 Absence of an Operational Test for Consciousness

Closely related is the absence of a reliable, substrate-independent test for consciousness. Existing measures in humans—subjective report, neural correlates, behavioral responsiveness, and clinical indices such as PCI—do not transfer to artificial systems. As a result, the claims in this paper rely on mechanistic and structural analysis rather than direct empirical detection. Developing an operational test for consciousness that applies across biological and artificial substrates remains an open challenge for the field.

7.3 Uncertainty About the Necessary and Sufficient Conditions for Consciousness

Although mainstream theories such as GWT, IIT, and AST propose mechanisms associated with consciousness, none provide a complete or universally accepted account. It remains unclear which mechanisms are necessary, which are sufficient, and whether consciousness requires specific biological substrates. This paper argues that current AI architectures do not implement the mechanisms proposed by these theories, but the theories themselves remain incomplete. Future developments in neuroscience or philosophy may refine or revise the mechanistic requirements for consciousness.

7.4 Architectural Novelty and Future Systems

The arguments presented here apply to contemporary AI architectures, including transformers, modular systems, and DMLG. It is possible that future architectures may differ in ways that challenge the assumptions of this paper. For example, systems with persistent internal state, self-modifying dynamics, or novel forms of recurrent integration may raise new questions about the relationship between computation and phenomenology. While the structural framework developed here provides strong reasons to doubt that such systems would be conscious, the possibility of radically new architectures cannot be excluded.

7.5 Interpretation of Self-Referential Behavior

This paper argues that self-referential language and introspective statements in AI systems are linguistic artifacts rather than indicators of subjective experience. While this interpretation is supported by mechanistic analysis and empirical anomalies such as hallucinations and inconsistent self-descriptions, it remains an interpretive stance. Alternative interpretations—for example, functionalist accounts that equate certain forms of self-modeling with minimal consciousness—are not addressed in detail here and remain topics of philosophical debate.

7.6 Limits of Canon-Based Arguments

The structural closure argument applied to DMLG relies on the assumption that the canonical graph fully constrains the system's behavior. While this assumption holds for current DMLG implementations, future extensions of the architecture may introduce new forms of learning, adaptation, or internal modeling. The extent to which canon-consistency can be preserved in more advanced versions of DMLG is an open question.

7.7 Open Questions

Several broader questions remain unresolved:

- What would constitute a substrate-independent test for consciousness?
- Are there minimal mechanistic conditions for consciousness that apply across biological and artificial systems?
- Could future architectures implement recurrent, self-updating processes that challenge the arguments presented here?
- To what extent can introspective language be decoupled from phenomenological claims in increasingly capable systems?
- How should the scientific community distinguish between functional competence, cognitive modeling, and subjective experience?

These questions highlight the conceptual and empirical gaps that remain in the study of consciousness. Addressing them will require interdisciplinary collaboration across neuroscience, philosophy, cognitive science, and artificial intelligence.

Despite these limitations, the arguments presented in this paper provide strong structural and mechanistic reasons to conclude that current AI systems are not conscious and that AGI does not require consciousness. The absence of a clear definition or operational test does not invalidate the structural separation between intelligence and phenomenology established in the IEGR framework.

8 Discussion

The arguments presented in this paper support a structural and mechanistic separation between intelligence, cognition, and consciousness. Building on the IEGR taxonomy, which defines AGI in purely functional and architectural terms, we have shown that none of the capabilities associated with general intelligence require subjective experience. Modern AI systems— including both transformer-based models and structurally distinct architectures such as the Dynamic Modular Language Graph (DMLG)—exhibit sophisticated cognitive behavior without implementing any of the mechanisms associated with consciousness in biological systems.

A central theme emerging from this analysis is that the appearance of conscious behavior in AI systems is best understood as a consequence of training data and representational structure, not as evidence of phenomenology. Large language models generate introspective or self-referential language because human language contains such patterns; DMLG generates similar patterns because its canonical graph encodes them. In both cases, the system produces linguistic artifacts that resemble conscious discourse without possessing an experiencing subject. This distinction is reinforced by the presence of hallucinations, inconsistent identities, and fabricated memories—behaviors incompatible with a stable phenomenological self but fully compatible with pattern continuation and narrative construction.

The comparison between transformers and DMLG further strengthens the argument. These architectures differ radically in their internal mechanisms: transformers rely on attention, context compression, and probabilistic sequence modeling, whereas DMLG relies on deterministic graph traversal, modular evaluators, and canon-consistent governance. Despite their architectural divergence, both systems exhibit similar limitations with respect to consciousness. Neither implements recurrent global broadcast, irreducible causal structure, self-updating internal models, or persistent temporal identity—mechanisms commonly associated with consciousness in theories such as GWT, IIT, and AST. The convergence of these independent lines of evidence suggests that consciousness is not an emergent property of architectural complexity, scale, or modularity.

This analysis also clarifies the relationship between AGI and consciousness. If AGI is defined structurally—as an IEGR-complete and IEGR-scaled system—then consciousness is neither a prerequisite nor a byproduct of general intelligence. AGI represents a phase shift in cognitive capacity, not in phenomenological status. Even ASI, if achieved through extensions of current architectures, would remain non-conscious unless it implemented mechanisms fundamentally different from those present in contemporary systems.

At the same time, the discussion highlights the conceptual and empirical challenges that remain in the study of consciousness. The absence of a universally accepted definition or operational test limits the ability to make definitive claims about artificial consciousness. Nevertheless, the structural and mechanistic arguments presented here provide a strong basis for concluding that current AI systems are not conscious and that consciousness is not required for AGI.

Overall, this paper contributes to a clearer conceptual landscape by disentangling intelligence, cognition, and consciousness, and by demonstrating that the capabilities of modern AI systems can be fully explained without invoking phenomenology. The implications extend beyond the specific architectures examined here: they suggest that future progress in AI should be evaluated in terms of functional competence and structural design, rather than assumptions about subjective experience.

Beyond its role as a counterexample to AI consciousness, DMLG serves as a conceptual enabler throughout the broader ecosystem. In the IEGR taxonomy, DMLG provides a concrete illustration of how general intelligence can be approached through structural and functional principles that differ fundamentally from transformers. In the context of minimal language competence, DMLG—when combined with a small LLM—was used to demonstrate how a micro-scale "sophisticated parrot" can achieve IEGR-completeness through the interaction of deterministic graph traversal and lightweight probabilistic modeling. This hybrid example highlights the distinction between linguistic behavior and underlying mechanisms. In this paper, DMLG functions as an orthogonal architecture whose deterministic, canon-consistent design makes explicit why consciousness does not arise from language generation, recursion, modularity, or learnable neurons. Its repeated appearance across these domains reflects not an approach toward consciousness, but its utility as a transparent, interpretable framework for understanding what modern AI systems are—and what they are not.

9 Conclusion

This paper has argued that intelligence, cognition, and consciousness are structurally distinct phenomena, and that contemporary artificial systems instantiate only the first two. Building on the IEGR taxonomy, we showed that AGI can be defined and achieved in purely functional and architectural terms, without invoking subjective experience. Modern AI systems—including transformers and the Dynamic Modular Language Graph (DMLG)—exhibit sophisticated cognitive behavior, generalization, and self-improving capabilities, yet none implement the mechanisms associated with consciousness in biological organisms.

The analysis proceeded along several independent lines. First, language generation, including introspective or self-referential language, was shown to be a computational artifact rather than a phenomenological signal. Second, the static and non-recurrent nature of current AI architectures precludes the dynamic, temporally extended processes required for subjective experience. Third, mainstream theories of consciousness (GWT, IIT, AST) do not predict consciousness in these systems, as none of the relevant mechanisms—global broadcast, irreducible causal structure, or self-updating internal models—are present. Fourth, hallucinations, inconsistent identities, and fabricated memories demonstrate that apparent introspection arises from narrative construction rather than from an experiencing subject. Finally, DMLG provided a structurally orthogonal argument: its canon-consistent architecture formally excludes the emergence of phenomenology, even in the presence of recursion, governance, and learnable neurons.

Taken together, these arguments support a unified conclusion: consciousness is neither a prerequisite for AGI nor a byproduct of current AI architectures. AGI represents a structural phase shift in cognitive capability, not a phenomenological transition. Even ASI, if achieved through extensions of existing architectures, would remain non-conscious unless it introduced mechanisms fundamentally different from those currently known.

The broader implication is that progress in artificial intelligence should be evaluated in terms of functional competence, structural design, and mechanistic transparency, rather than assumptions about subjective experience. While the scientific study of consciousness remains conceptually and empirically unresolved, the evidence presented here provides strong grounds for separating intelligence from phenomenology in artificial systems. Consciousness, as far as current theory and mechanism indicate, is not produced by the architectures that underlie modern AI—and is not required for the development of general or even superintelligent artificial agents.

A Reflective Note on Consciousness, Triadic Structure, and Artificial Systems

This appendix offers a conceptual reflection situated within the broader Triadic Cosmos ecosystem. It is not part of the main argument of the paper and does not introduce new claims about physics, cosmology, or artificial intelligence. Instead, it highlights an internal observation that follows naturally from the structural alignment between the Triadic Cycle and the IEGR model of intelligence.

The Triadic Cycle describes the universe in terms of three irreducible structural roles—energy (E), geometry (G), and information (I)—linked through the recursive update operator R that generates emergent temporal structure. In the IEGR taxonomy, general intelligence is characterized by the same triadic roles, supplemented by recursion as a governing operator. This structural isomorphism does not imply that the universe is intelligent, nor that intelligence is a cosmological property. Rather, it shows that both physical law and intelligent behavior can be described using the same minimal architectural grammar.

Within this framework, consciousness does not appear as a triadic role, nor as a consequence of triadic recursion. The universe, as described by the Triadic Cycle, reorganizes energetic, geometric, and informational structure, but it does not instantiate phenomenological experience or an experiencing subject. The same is true for IEGR: intelligence is defined structurally and functionally, without reference to subjective awareness. Consciousness is therefore not excluded by the triadic architecture, but it is not required or generated by it.

This observation has a modest but relevant implication for artificial systems. If both the universe and intelligent behavior can be described triadically, and if consciousness is not a triadic role, then artificial systems that instantiate triadic structure—whether through transformers, DMLG, or other architectures—do not acquire consciousness merely by participating in the same structural grammar. Consciousness is not a direct consequence of triadic organization, emergent time, or recursive dynamics.

This reflection does not claim that consciousness is impossible in artificial systems, nor that biological consciousness is unrelated to physical law. It simply notes that within the Triadic Cosmos framework, consciousness is not a structural component of the triadic architecture that governs both physical systems and artificial intelligence. As such, the absence of consciousness in current AI systems is consistent with the broader structural ontology of the ecosystem.

This appendix is offered as an interpretive remark for readers familiar with the Triadic Cosmos program. It is not intended as a universal claim about consciousness, but as a clarification of how the present paper fits within the conceptual landscape of the ecosystem.

B Triadic Cosmos Ecosystem Context

This paper is part of the *Triadic Cosmos* ecosystem, which develops a unified conceptual and computational framework spanning physical ontology and modular architectures for intelligent systems. All materials, extended notes, and evolving documentation are available in the public repository: `https://github.com/triadic-cosmos`.

References

Baars, Bernard J. (1988). *A Cognitive Theory of Consciousness*. Cambridge University Press.

Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio (2015). "Neural Machine Translation by Jointly Learning to Align and Translate". In: *International Conference on Learning Representations (ICLR)*.

Bubeck, Sébastien et al. (2023). "Sparks of Artificial General Intelligence: Early Experiments with GPT-4". In: *arXiv preprint arXiv:2303.12712*.

De Witte, Joachim (2026a). "A Dynamic and Modular Architecture for Language Modeling". In: *arXiv preprint*.

— (2026b). "A Structural Taxonomy of Intelligence: IEGR-Scale, Parrot Regimes, and the Conditions for Human-Like AGI". In: *arXiv preprint*.

— (2026c). "Language as Graph Traversal: The DMLG Architecture". In: *arXiv preprint*.

Graziano, Michael S. A. and Sabine Kastner (2013). "Human Consciousness and Its Relationship to Social Neuroscience: A Novel Hypothesis". In: *Cognitive Neuroscience* 4.3–4, pp. 113–120.

LeCun, Yann, Yoshua Bengio, and Geoffrey Hinton (2015). "Deep Learning". In: *Nature* 521, pp. 436–444.

Tononi, Giulio (2004). "An Information Integration Theory of Consciousness". In: *BMC Neuroscience* 5.42.

Turing, Alan M. (1950). "Computing Machinery and Intelligence". In: *Mind* 59.236, pp. 433–460.