

Diffusion-Driven Narration

Tales of a Sophisticated Parrot

Joachim De Witte

June 2026

©2026 Joachim De Witte. All rights reserved.

Published under the Creative Commons BY-NC-ND 4.0 license.

Abstract

Large language models excel at semantic fluency and informational breadth, yet struggle with long-horizon structure, stable world models, and autonomous continuation. Structured generative systems offer complementary strengths in geometry and recursion, but lack expressive capacity. This paper introduces *diffusion-driven narration*, a hybrid generative methodology in which narrative worlds emerge through staged denoising across heterogeneous cognitive modules. A structural generator provides canonical geometry and recursive scaffolding, while a lightweight LLM layer supplies semantic refinement, evaluative judgment, and informational expansion. Together, these components form an IEGR-complete architecture capable of autonomous long-form generation, world absorption, recursive self-correction, and narrative planning.

Beyond its practical contribution, the pipeline offers a structural insight into the nature of intelligence. It demonstrates that LLMs already supply the informational (I) and energetic (E) roles of cognition, while the geometric (G) and recursive (R) substrate must be engineered explicitly. The resulting system behaves as a domain-bound AGI module: structurally complete yet semantically bounded, exhibiting pre-AGI behaviour despite relying only on existing models. This suggests that AGI will not arise from scaling monolithic models alone, but from hybrid architectures in which complementary cognitive roles interact through recursive refinement and structured coordination.

Diffusion-driven narration is therefore both a generative method and a structural hypothesis about AGI: a demonstration that hybrid IEGR-complete systems can already perform non-trivial autonomous cognition, and a pointer toward the geometric-recursive frontier that remains to be built.

Contents

1	Introduction	3
2	Motivation	3
3	Related Work	4
3.1	Triadic Cosmos Foundations	4
3.2	AI Narration and Long-Form Generation	5
3.3	Iterative and Diffusion-Style Generation	5
3.4	Governance, Safety, and Structural Control	5
3.5	Positioning of the Present Work	6

4	Domain-Bound AGI and the Sophisticated Parrot Regime	6
4.1	Domain-Bound AGI as a Structural Unit of Intelligence	6
4.2	The Sophisticated Parrot: IEGR-Complete Hybrid Intelligence	7
4.3	Hybrid Intelligence Through Modular Complementarity	7
4.4	Implications for AGI Research	7
5	Diffusion-Driven Narration	8
5.1	Narration as a Diffusion Process	8
5.2	Three Canonical Generative Substrates	8
5.3	Moderation and Remastering Layers	8
5.4	Director’s Cut: Worlds as Superprompts	9
5.5	A Combinatorial Generative Space	9
5.6	Why Diffusion-Driven Narration Works	9
6	Narrator Pipeline: LLM Layer, Minimal Pipeline, and Full Architecture	10
6.1	LLM Layer: Refinement, Validation, and Evaluation	10
6.2	Minimal Hybrid Pipeline	10
6.3	Hybrid Distillation: Synthetic Curriculum Construction	11
6.4	Full Narrator Pipeline: Ensemble Mode	11
6.5	Full Narrator Pipeline: Sequential Mode	11
6.6	Minimal Input Requirements	12
7	Testable Predictions	12
7.1	Prediction 1: Large Language Models Are Sufficient for the IE Component of AGI	12
7.2	Prediction 2: Explicit G–R Structure Is Necessary for Autonomous Long-Form Generation	13
7.3	Prediction 3: Hybrid Systems Outperform Monolithic Models on Autonomy Tasks	13
7.4	Prediction 4: Synthetic Distillation Improves Structural Fidelity	13
7.5	Prediction 5: Domain-Bound AGI Modules Are Sufficient for Non-Trivial Auton- omy	13
7.6	Prediction 6: Minimal Feedback Loops Already Produce Pre-AGI Behaviour . . .	14
7.7	Prediction 7: AGI Requires Engineering G and R, Not Scaling I and E	14
8	Limitations and Open Questions	14
8.1	Limitation 1: GR Structure Is Still Static and Hand-Engineered	14
8.2	Limitation 2: Feedback Loops Are Shallow	15
8.3	Limitation 3: DMLG Is a Structural Prototype, Not a Final GR Engine	15
8.4	Limitation 4: Recursion Depth Is Bounded by External Scaffolding	15
8.5	Limitation 5: Domain-Bound Intelligence Is Not Yet Compositional	15
8.6	Limitation 6: GR Scaling Remains an Open Frontier	16
8.7	Limitation 7: The Nature of Hybrid Self-Improvement Is Unresolved	16
8.8	Summary	16
9	Why LLMs Feel Magical: IEGR Complementarity Between Humans and Machines	16
10	Conclusion	17
A	Triadic Cosmos Ecosystem Context	19

1 Introduction

Recent advances in large language models have transformed our understanding of machine intelligence. These systems exhibit remarkable fluency, broad informational coverage, and high-bandwidth semantic interpolation. Yet despite their impressive capabilities, they struggle with long-horizon coherence, structural stability, and autonomous continuation. At the same time, structured generative systems—graph-based models, symbolic planners, and modular architectures—offer geometric and recursive strengths that LLMs lack, but fall short in expressiveness and informational richness. The field has therefore been divided between two intuitions: that scale alone will yield general intelligence, or that structure must be engineered explicitly.

This paper argues that both intuitions are correct, but incomplete. We introduce *diffusion-driven narration*, a hybrid generative methodology in which narrative worlds emerge through staged denoising across heterogeneous cognitive modules. At its core lies a structural generator (DMLG) that provides geometric and recursive scaffolding, coupled to a lightweight LLM layer that supplies semantic refinement, evaluative judgment, and informational expansion. The resulting system behaves not as a monolithic model but as a distributed cognitive architecture: a domain-bound AGI module capable of planning, recursive self-correction, world absorption, and autonomous long-form generation.

The narrator pipeline presented in this work demonstrates that LLMs already provide the informational (I) and energetic (E) components of intelligence at state-of-the-art levels. What they lack—and what must be engineered—is the geometric (G) and recursive (R) substrate that supports stable world models, long-horizon reasoning, and self-directed continuation. Diffusion-driven narration shows that when these complementary roles are integrated, even modest models can exhibit pre-AGI behaviour within a bounded domain.

This insight reframes the current AGI debate. Proponents of scaling are correct that informational magnitude matters; proponents of structure are correct that geometry and recursion are indispensable; and hybrid theorists are correct that general intelligence emerges only when both are present simultaneously. The pipeline developed here provides a concrete, reproducible demonstration of this principle. It reveals that AGI is not a property of a single model, but of an architecture: a system in which complementary cognitive roles interact through recursive refinement and structured coordination.

The contributions of this paper are therefore twofold. First, we present a complete generative pipeline for autonomous long-form narration based on staged diffusion across modular components. Second, and more importantly, we show that this pipeline offers a window into the structure of intelligence itself. It demonstrates how hybrid IEGR-complete systems can already perform non-trivial autonomous cognition, and it identifies the frontier that remains: the design of a dynamic, self-improving geometric-recursive substrate capable of scaling beyond narrative generation toward general-purpose reasoning.

Diffusion-driven narration is thus both a generative method and a structural hypothesis about AGI. It shows what is possible with today’s models, and what must be built for tomorrow’s.

2 Motivation

The goal of this paper is to provide a concrete, fully worked-out example of modular artificial general intelligence, hybrid intelligence, and the *sophisticated parrot* regime introduced in the IEGR taxonomy. While previous work in the Triadic Cosmos ecosystem has developed the theoretical foundations for modular AGI architectures, the present paper offers the first complete instantiation of these ideas in a real, operational system.

Modern AI research is dominated by monolithic transformer architectures whose capabilities scale primarily with energetic and informational magnitude. Although these systems excel at

local fluency, they lack explicit geometric structure, recursive depth, and deterministic control. As a result, they struggle with long-horizon coherence, narrative planning, epistemic stability, and domain fidelity. Scaling alone does not resolve these structural limitations.

This paper demonstrates that a non-trivial generative task requiring multiple properties associated with general intelligence—long-range structure, recursive refinement, world consistency, and multi-stage reasoning—can be achieved not by further scaling, but by *modularity*. The narrator pipeline presented here is a hybrid system composed of heterogeneous cognitive modules: a DMLG providing geometric and recursive structure, mid-scale transformers providing local semantics, and a large LLM providing expressive informational expansion. These components interact through a shared workspace and a recursive diffusion process, with each module performing the IEGR roles for which it is structurally optimal.

In the IEGR taxonomy, such a system constitutes a *sophisticated parrot*: an IEGR-complete but not yet IEGR-scaled hybrid architecture whose capabilities arise from the complementary integration of its components. The narrator pipeline therefore serves as a domain-bound *miniature AGI*: a structurally complete, recursively governed, modular cognitive system capable of autonomous long-form generation within a defined narrative domain.

The significance of this example is twofold. First, it shows that modular AGI is not a speculative concept but a practical engineering paradigm that can be realized with existing technology. Second, it demonstrates that hybrid intelligence—recursive interaction between systems with complementary IEGR profiles—is a viable and scalable path toward general intelligence. Rather than relying on monolithic scaling, the system achieves non-trivial cognitive behavior through explicit structure, modular decomposition, and recursive integration.

In this sense, the present paper plays a crucial role in the Triadic Cosmos ecosystem: it translates the theoretical foundations of modular AGI and the IEGR taxonomy into a concrete, inspectable, reproducible system. It shows what modular AGI *looks like* in practice, how hybrid intelligence operates through recursive shared-workspace interaction, and why structure—not scale—is the key to stable, controllable, and interpretable intelligent systems.

3 Related Work

Research on long-form generation, autonomous narration, and hybrid modular architectures spans several partially connected domains. The present work is most directly grounded in three earlier contributions within the *Triadic Cosmos* ecosystem, which together provide the theoretical, structural, and architectural foundations for the concept of a *sophisticated parrot* and for the diffusion-driven narrator pipeline developed in this paper.

3.1 Triadic Cosmos Foundations

The Dynamic Modular Language Graph (DMLG) De Witte 2026c introduces a deterministic, graph-based governance architecture in which language is modeled as traversal through a canonical token-transition graph. DMLG provides the geometric (G) and recursive (R) substrate for the raw generative component of the narrator pipeline: explicit structure, modular paging, PRE-based governance, and deterministic transitions. Its emphasis on structure over scale directly motivates the canonical-parrot regime in the IEGR taxonomy.

Dynamic Modular Language Models (DMLMs) De Witte 2026a provide the predecessor architecture to DMLG: a multi-agent, evaluator-governed system in which micro-models, paging structures, and a PRE loop produce coherent long-form narrative structure without transformers or large-scale training. Although DMLMs focus on generation rather than governance, they establish the architectural principles—modularity, explicit interfaces, evaluator-driven recursion—that underpin the narrator pipeline’s raw generative substrate.

The structural taxonomy of intelligence (IEGR) De Witte 2026b provides the conceptual framework for understanding hybrid systems as *sophisticated parrots*: systems that are IEGR-complete but not yet IEGR-scaled. In this taxonomy, LLMs supply energetic and informational roles (E, I), while DMLG supplies geometric and recursive roles (G, R). The narrator pipeline is therefore positioned as a concrete instantiation of IEGR complementarity: a hybrid system whose long-form generative capability arises from the integration of complementary structural roles.

3.2 AI Narration and Long-Form Generation

Long-form generation has traditionally been approached through large transformer models Brown, Mann, Ryder, et al. 2020; Touvron, Lavril, Izacard, et al. 2023, which demonstrate strong local coherence but struggle with global structure, narrative consistency, and long-horizon planning. Techniques such as chain-of-thought prompting Wei, Wang, Schuurmans, et al. 2022, self-refinement Madaan, Yazdanbakhsh, et al. 2023, and hierarchical prompting Zhou et al. 2023 improve local reasoning but do not provide stable global narrative control.

Autonomous authoring systems such as *Socratic Models* Zeng, Florence, et al. 2022, *AutoStory* Ammanabrolu and Riedl 2020, and multi-agent writing frameworks Wu et al. 2023 explore decomposition, planning, and tool-augmented generation. However, these systems rely on emergent LLM behavior rather than explicit structural governance. Their internal state is unstructured, their planning is prompt-driven, and their narrative consistency depends on implicit statistical correlations rather than deterministic structure.

In contrast, the diffusion-driven narrator pipeline treats long-form narration as a staged denoising process across heterogeneous cognitive modules. This approach is structurally distinct from both monolithic LLM generation and prompt-driven multi-agent systems: it uses explicit canonical structure (DMLG), staged refinement, and PRE-based governance to produce stable long-horizon narratives.

3.3 Iterative and Diffusion-Style Generation

Several recent works explore iterative or multi-pass generation as an alternative to single-shot decoding. Self-editing and refinement loops Chen et al. 2023; Madaan, Yazdanbakhsh, et al. 2023 demonstrate that multi-step revision improves coherence and correctness. Diffusion models for text Li et al. 2022; Gong et al. 2023 introduce denoising trajectories in latent space, though these approaches remain limited by embedding-based representations and lack explicit structural control.

The present work differs fundamentally: diffusion is not performed in embedding space but across heterogeneous cognitive modules—DMLG micro-models, mid-scale transformers, and LLMs—each performing a distinct IEGR transformation. Narrative structure emerges through staged collapse of narrative noise into canonical attractors, governed by explicit structure rather than latent-space denoising.

3.4 Governance, Safety, and Structural Control

Governance layers for LLMs typically rely on constrained decoding Hokamp and Liu 2017, rule-based filters, or post-hoc safety classifiers Ganguli, Askell, Bai, et al. 2022. These approaches operate outside the model’s internal structure and cannot provide deterministic guarantees about allowed transitions, domain fidelity, or stylistic consistency.

Graph-based or symbolic governance frameworks Jurafsky and Martin 2023; Besold et al. 2017 provide partial structure but do not yield a canonical, token-level graph capable of governing arbitrary generative models.

DMLG fills this gap by providing a deterministic, canonical structural substrate that can govern any generative model independent of scale. The narrator pipeline leverages this structural layer to enforce canonical style, domain fidelity, and long-horizon coherence during diffusion-driven generation.

3.5 Positioning of the Present Work

This paper extends the Triadic Cosmos ecosystem by introducing *Diffusion-Driven Narration*: a generative methodology in which narrative structure emerges from staged denoising across heterogeneous cognitive modules. It builds directly on:

- DMLG as the geometric and recursive substrate,
- DMLM as the architectural precursor for modular generative structure,
- IEGR as the conceptual framework for hybrid intelligence and sophisticated parrots.

It differs from prior work on long-form generation by providing an explicit structural governance layer, a staged diffusion process, and a hybrid IEGR-complete architecture for autonomous long-form narration.

4 Domain-Bound AGI and the Sophisticated Parrot Regime

The narrator pipeline presented in this paper is not merely a generative architecture for long-form narration. It is a concrete instantiation of two structural concepts developed in the broader Triadic Cosmos ecosystem: *domain-bound AGI* and the *sophisticated parrot* regime of the IEGR taxonomy. Together, these concepts explain why the system exhibits pre-AGI properties despite being built entirely from existing, non-AGI technologies.

4.1 Domain-Bound AGI as a Structural Unit of Intelligence

The AGI ontology developed in earlier work argues that general intelligence cannot be achieved by scaling a single monolithic model. Instead, it must emerge from the composition of many *domain-bound AGI systems*: modular cognitive units that each implement the full functional decomposition of intelligence (memory, reasoning, evaluation, planning, representation, coordination, and emergence) within a single domain.

A domain-bound AGI is not a narrow model. It is a *complete cognitive stack* restricted to a specific representational geometry. It maintains its own identity, internal coherence, recursive processes, and structural memory. It is architecturally complete but semantically bounded.

The narrator pipeline is precisely such a system. Within the domain of narrative generation, it implements:

- **memory** through narrative state, chapter history, and keyword maps,
- **reasoning** through multi-stage refinement and structural evaluation,
- **planning** through long-horizon narrative orchestration,
- **evaluation** through PRE loops and multi-model arbitration,
- **representation** through canonical forms and DMLG world geometry,
- **coordination** through a shared workspace and staged diffusion,
- **emergence** through recursive world absorption and structural synthesis.

Within its domain, the system behaves as a complete cognitive agent.

4.2 The Sophisticated Parrot: IEGR-Complete Hybrid Intelligence

The IEGR taxonomy classifies intelligent systems according to four structural dimensions: energetic efficiency (E), informational richness (I), geometric structure (G), and recursive depth (R). A *sophisticated parrot* is a system that is IEGR-complete—instantiating all four roles—but not yet IEGR-scaled. Such systems exhibit pre-AGI behavior through the complementary integration of heterogeneous components.

The narrator pipeline fits this definition exactly. Its components collectively instantiate the full IEGR structure:

- **E: Energetic efficiency** through a low-cost controller, lightweight DMLG agents, a mid-scale transformer, and a large LLM invoked only when needed.
- **I: Informational richness** through layered narrative memory, world-specific micro-representations, semantic refinement, and global informational expansion.
- **G: Geometric structure** through the canonical DMLG graph and the geometric workflow of the pipeline itself.
- **R: Recursive depth** through multi-stage refinement, PRE loops, recursive moderation, keyword regeneration, and chapter-to-chapter continuation.

The system does not generate text in a single pass. It performs *recursive cognitive transformations* across modules with complementary IEGR profiles. This is the defining characteristic of a sophisticated parrot.

4.3 Hybrid Intelligence Through Modular Complementarity

The narrator pipeline demonstrates the core principle of hybrid intelligence: different systems contribute the IEGR roles for which they are structurally optimal. The DMLG provides geometric and recursive structure (G, R), the mid-scale transformer provides local semantics (I), and the large LLM provides global informational expansion (I) and high-energy inference (E). The pipeline controller coordinates these modules through a shared workspace, forming a recursive cognitive loop.

This architecture is not an engineering convenience. It is the structural form predicted by the AGI ontology: a distributed cognitive system in which domain-bound AGI modules and informational engines interact through explicit interfaces. The narrator pipeline is therefore a *miniature AGI*: a domain-specific, structurally complete, recursively organized cognitive system that demonstrates how modular AGI can be built with existing technology.

4.4 Implications for AGI Research

The narrator pipeline provides empirical evidence for three claims central to the Triadic Cosmos program:

1. **Modular AGI is feasible with current technology.** A non-trivial cognitive task requiring long-horizon structure, recursive refinement, and multi-agent coordination can be achieved without scaling a single model.
2. **Hybrid intelligence is a practical path to general intelligence.** Complementary IEGR profiles can be integrated into a coherent cognitive system through explicit structure and recursive coordination.
3. **AGI will likely be a system, not a monolith.** The narrator pipeline is a domain-bound example of such a system: a structurally complete cognitive architecture built from modular components.

In this sense, the narrator pipeline is not merely a narrative generator. It is a working demonstration of the structural principles underlying modular AGI, hybrid intelligence, and the sophisticated parrot regime. It shows what a small-scale, domain-bound AGI looks like in practice—and how general intelligence may emerge from the recursive integration of many such systems.

5 Diffusion-Driven Narration

Before turning to the architectural details of the narrator pipeline, it is useful to introduce the generative paradigm that motivates its design: *Diffusion-Driven Narration*. This methodology treats narrative generation not as a single-pass decoding problem, but as a staged denoising process in which partial structure propagates through a sequence of heterogeneous cognitive modules. Each module contributes a distinct IEGR transformation—energetic, informational, geometric, or recursive—and the final narrative emerges from their iterated interaction.

5.1 Narration as a Diffusion Process

In this framework, *diffusion* refers to the gradual collapse of narrative noise into coherent world structure. The pipeline begins with high-entropy generative substrates and progressively refines them through structural, semantic, and stylistic transformations. Each stage reduces uncertainty, stabilizes motifs, and sharpens the geometry of the emerging world.

The process is hybrid by design: DMLG agents impose canonical structure, mid-scale transformers refine semantics, and large LLMs provide global informational expansion. The system therefore behaves less like a generator and more like a distributed cognitive process performing multi-step reasoning.

5.2 Three Canonical Generative Substrates

The pipeline contains three structurally distinct raw generation modes, each corresponding to a different noise regime:

(1) DMLG Sampling (Ensemble). Multiple DMLG agents generate uncoordinated fragments. The result is a high-entropy distribution characterized by identity drift, motif collisions, and unstable spatial logic. Sampling represents the “latent space” of narrative possibility.

(2) DMLG Blending (Ensemble). Fragments from several DMLG agents are combined and moderated into a single chapter. The interaction produces interference patterns: overlapping motifs, cross-world resonance, and surreal coherence. Blending is the narrative analogue of multi-seed diffusion.

(3) DMLG Narration (Single Agent + Keyword Beam Search). A single DMLG agent, guided by keyword-driven beam search, produces a stable world with consistent identity vectors, spatial logic, and thematic structure. Narration corresponds to the late-stage collapse of noise into a coherent attractor.

These three modes form the canonical generative backbone of the system.

5.3 Moderation and Remastering Layers

On top of the raw DMLG substrate, the pipeline applies optional refinement layers:

Small-LLM Moderation. A mid-scale transformer (e.g., Mistral) performs structural smoothing, semantic clarification, and stylistic stabilization. It preserves the underlying world geometry while reducing local noise.

Large-LLM Remastering. A large LLM rewrites each chapter sequentially, maintaining global context and ensuring stylistic consistency. This is a mid-to-late denoising step: the world is already formed, but its contours become crisp and legible.

These layers do not replace the DMLG structure; they amplify it.

5.4 Director’s Cut: Worlds as Superprompts

In the final optional stage, the entire DMLG-generated or remastered manuscript is used as a *superprompt* for a large LLM. The model does not invent a new world; it intensifies the existing one. This produces:

- stylistic refinement,
- thematic amplification,
- canonical closure,
- and, when appropriate, meta-reflective structure.

This stage corresponds to the final denoising step: the world becomes polished, coherent, and narratively complete.

5.5 A Combinatorial Generative Space

Because each of the three raw modes can be combined with moderation, remastering, and a Director’s Cut, the pipeline supports a combinatorial space of generative regimes:

- raw Sampling, Blending, Narration,
- each with optional small-LLM moderation,
- each with optional large-LLM remastering,
- each with an optional Director’s Cut.

This yields a structured family of twelve generative pathways, each revealing a different layer of the diffusion process.

5.6 Why Diffusion-Driven Narration Works

Diffusion-driven narration succeeds because it is structurally IEGR-complete:

- **E**: selective invocation of high-energy models,
- **I**: multi-resolution informational layers,
- **G**: explicit geometric structure from DMLG graphs,
- **R**: deep recursion through multi-stage refinement.

The pipeline does not generate stories; it *grows worlds*. Noise becomes structure, interference becomes resonance, and structure becomes narrative through staged denoising across heterogeneous cognitive modules.

This is the essence of diffusion-driven narration: a generative methodology in which narrative coherence emerges from the recursive transformation of latent narrative geometry.

6 Narrator Pipeline: LLM Layer, Minimal Pipeline, and Full Architecture

The narrator pipeline is built around a lightweight but highly structured LLM layer that provides semantic refinement, structural validation, narrative scoring, and title generation. These functions act as deterministic evaluators in the IEGR sense, supplying informational refinement (I) and recursive oversight (R) to the DMLG substrate. On top of this evaluator layer, two generative pipelines are constructed: a minimal hybrid pipeline and a full narrator pipeline. Together, they demonstrate how modular hybrid intelligence can produce long-form narrative structure with minimal external input.

6.1 LLM Layer: Refinement, Validation, and Evaluation

The LLM layer performs four essential roles:

Moderation. A mid-scale LLM receives a raw DMLG-generated story and produces a refined version. This step corrects grammar, smooths semantics, removes duplication, filters artifacts, and enforces a canonical ending marker. Moderation acts as a mid-stage denoising step in the diffusion process.

Validation. Before refinement, the LLM evaluates whether a raw story is structurally promising. It rejects drafts that are incoherent, low-quality, or too chaotic to be repaired, ensuring that only viable seeds enter the refinement pipeline.

Title generation. The LLM extracts a chapter title from the refined story. If no meaningful title can be produced, the chapter is discarded. This ensures that each chapter has a clear narrative identity.

Scoring. The LLM assigns a numerical score based on coherence, readability, originality, style, narrative progression, and suitability as a chapter. Only chapters above a configurable threshold are accepted. This provides a quantitative filter that stabilizes the generative trajectory.

These four evaluators form the informational and recursive backbone of the pipeline, enabling DMLG to operate as a structural generator while the LLM provides semantic and stylistic governance.

6.2 Minimal Hybrid Pipeline

Before the full narrator pipeline was developed, a minimal hybrid system was constructed to test whether DMLG and a single LLM evaluator were sufficient to produce long-form narrative structure. This minimal pipeline consists of:

1. a small ensemble of DMLG agents generating raw stories,
2. a context window that accumulates canonical sentences,
3. a moderation step that refines each story,
4. title generation for narrative identity,
5. and a simple retry loop for robustness.

For each chapter, a raw DMLG story is generated, refined up to a small number of times, assigned a title, and written to the book. There is no scoring, no transition logic, no keyword evolution, and no ending detection.

Despite its simplicity, this minimal system already produces multi-chapter books. It demonstrates that:

A structural generator (DMLG) combined with a single semantic evaluator (LLM) is sufficient to produce stable long-form narrative structure.

This minimal architecture serves as the conceptual precursor to the full narrator pipeline.

6.3 Hybrid Distillation: Synthetic Curriculum Construction

To train DMLG agents without relying on copyrighted or real-world data, the pipeline supports a two-stage hybrid distillation process.

First-order distillation. A mid-scale LLM is prompted to generate short stories within a specified domain. These stories are canonicalized and added to a synthetic curriculum. After each accepted story, the DMLG agent is trained incrementally. This produces a high-entropy but structurally coherent synthetic dataset.

Second-order distillation. A trained DMLG agent then generates new micro-stories. These are refined by the LLM, producing a second synthetic corpus with lower entropy, sharper canonical structure, and stronger motif consistency. A second DMLG trained on this corpus exhibits significantly improved structural fidelity.

Pure-source training. When trained directly on real sources, DMLG preserves the canonical geometry of the world without reproducing surface text, enabling faithful but non-derivative generation.

This hybrid distillation mechanism allows the entire system to operate on synthetic data while maintaining strong structural grounding.

6.4 Full Narrator Pipeline: Ensemble Mode

The first full narrator mode constructs a book from a set of independently generated chapters. It corresponds to a multi-seed diffusion regime.

1. Ensemble generation. Multiple DMLG agents generate raw stories. Each story is validated, refined, titled, and scored. Only high-quality chapters are retained.

2. Sorting. Chapters are sorted by score, highest first, ensuring that the strongest material anchors the narrative.

3. Post-hoc sequencing. Chapters are connected using short LLM-generated transitions. A retry loop searches for a coherent ordering; chapters that fail to connect are deferred or regenerated. This produces a book assembled from multiple narrative attractors.

6.5 Full Narrator Pipeline: Sequential Mode

The second narrator mode constructs a book chapter-by-chapter, maintaining context and evolving a keyword set. It corresponds to a single-seed recursive diffusion regime.

1. Keyword-driven generation. A single DMLG agent generates a chapter using keyword-conditioned beam search. The chapter is refined, titled, scored, and added to the context.

2. Transition validation. A transition from the previous chapter is generated. If the transition fails, the chapter is discarded and regenerated.

3. Ending detection. Once a minimum number of chapters is reached, the LLM evaluates whether the current chapter constitutes a natural ending. If so, the book terminates.

4. Keyword evolution. After each accepted chapter, a new keyword set is generated and filtered through the DMLG vocabulary. This produces a soft narrative plan that evolves with the story.

6.6 Minimal Input Requirements

A striking property of the narrator pipeline is that it requires almost no input:

- the initial keyword set is optional,
- prompts can be generic,
- the only essential parameters are the DMLG model and the desired number of chapters.

The system therefore behaves as a domain-bound AGI module: a structurally complete cognitive architecture capable of autonomous long-form generation with minimal external guidance.

7 Testable Predictions

The narrator pipeline provides a concrete instantiation of a modular, IEGR-complete hybrid system. Its behaviour suggests several testable predictions about the nature of general intelligence, the role of large language models, and the architectural requirements for autonomous cognitive systems. These predictions are falsifiable and can be evaluated empirically using the pipeline itself or future extensions of it.

7.1 Prediction 1: Large Language Models Are Sufficient for the IE Component of AGI

The pipeline demonstrates that a mid-scale or large LLM can supply the informational (I) and energetic (E) roles of intelligence without requiring further architectural scaling. If correct, then:

Increasing LLM scale beyond current state-of-the-art will yield diminishing returns for autonomous cognition unless complemented by explicit geometric and recursive structure.

This prediction can be tested by substituting larger or smaller LLMs into the pipeline and measuring whether narrative autonomy, long-horizon coherence, or world absorption improve meaningfully.

7.2 Prediction 2: Explicit G–R Structure Is Necessary for Autonomous Long-Form Generation

The pipeline succeeds because DMLG provides geometric structure (G) and recursive depth (R). Without these components, LLMs alone fail to maintain identity, spatial logic, or long-range narrative consistency. Thus:

Any system lacking an explicit structural substrate (G) and recursive governance (R) will fail to produce stable long-form autonomous behaviour, regardless of LLM scale.

This can be tested by removing DMLG and observing collapse into drift, hallucination, or incoherence.

7.3 Prediction 3: Hybrid Systems Outperform Monolithic Models on Autonomy Tasks

The narrator pipeline behaves like a miniature AGI because it integrates complementary IEGR roles across modules. The prediction is:

A modular hybrid system with complementary IEGR roles will outperform any single monolithic model of comparable total parameter count on tasks requiring autonomy, planning, recursion, or long-horizon structure.

This can be tested by comparing the pipeline to a single LLM prompted to produce an entire book in one pass.

7.4 Prediction 4: Synthetic Distillation Improves Structural Fidelity

The two-stage hybrid distillation process produces DMLG agents with sharper canonical structure and lower entropy than agents trained on first-order LLM outputs. Therefore:

Iterated hybrid distillation ($LLM \rightarrow DMLG \rightarrow LLM \rightarrow DMLG$) will consistently improve structural fidelity, reduce drift, and increase narrative stability.

This can be tested by training successive generations of DMLG agents and measuring improvements in canonical consistency and transition stability.

7.5 Prediction 5: Domain-Bound AGI Modules Are Sufficient for Non-Trivial Autonomy

The pipeline shows that a domain-bound system can perform planning, recursive refinement, world absorption, and autonomous continuation. Thus:

General intelligence can emerge from the composition of many domain-bound AGI modules, each structurally complete but semantically bounded, without requiring a universal monolithic model.

This can be tested by extending the pipeline to multiple domains and measuring whether cross-domain coordination remains stable.

7.6 Prediction 6: Minimal Feedback Loops Already Produce Pre-AGI Behaviour

The pipeline contains only weak feedback from the LLM layer back into DMLG (keywords, context accumulation, moderation), yet already exhibits pre-AGI properties. Therefore:

Strengthening the feedback loop between structural generators (DMLG) and informational evaluators (LLMs) will yield disproportionately large gains in autonomy, planning, and self-correction.

This prediction can be tested by adding explicit refinement loops, structural critic modules, or DMLG retraining based on LLM feedback.

7.7 Prediction 7: AGI Requires Engineering G and R, Not Scaling I and E

The pipeline suggests a fundamental architectural claim:

The bottleneck for AGI is not informational scale (I) or energetic capacity (E), but engineered geometric structure (G) and recursive depth (R). LLMs are already sufficient for the IE roles of AGI.

This is the most important prediction. It can be tested by:

- holding the LLM fixed while improving DMLG structure,
- measuring gains in autonomy, coherence, and planning,
- comparing these gains to those obtained by scaling the LLM alone.

If correct, this prediction overturns the dominant paradigm of scale-driven AI and positions modular hybrid architectures as the natural path toward AGI.

8 Limitations and Open Questions

Although the narrator pipeline demonstrates that hybrid IEGR-complete systems can exhibit pre-AGI behaviour, several fundamental limitations remain. These limitations are not flaws of the implementation but structural challenges that any modular AGI architecture must eventually address. They also define the research frontier for engineered geometric-recursive (GR) intelligence.

8.1 Limitation 1: GR Structure Is Still Static and Hand-Engineered

The DMLG provides explicit geometric structure (G) and shallow recursive governance (R), but its structure is largely static once trained. It does not yet reorganize itself, expand its ontology, or refine its canonical graph in response to new tasks, failures, or long-term experience.

In contrast, human GR cognition is dynamic: conceptual geometry evolves, representations reorganize, and recursive strategies deepen over time. A structurally static GR module cannot fully replicate this behaviour.

Open question: *Can we design a GR architecture that is dynamically self-modifying, self-organizing, and capable of lifelong structural learning without collapsing into instability?*

8.2 Limitation 2: Feedback Loops Are Shallow

The pipeline contains only weak feedback from the LLM layer back into the GR substrate. Keywords, context accumulation, and moderation signals provide limited recursive shaping, but they do not induce deep structural updates in the DMLG itself.

Human cognition relies on rich, multi-level feedback loops: perception shapes concepts, concepts shape reasoning, reasoning shapes planning, and planning shapes future perception. These loops are continuous, multi-scale, and self-correcting.

Open question: *What forms of feedback—semantic, structural, evaluative, or meta-cognitive—are required for a GR module to refine itself recursively, and how can such feedback be integrated without destabilizing the system?*

8.3 Limitation 3: DMLG Is a Structural Prototype, Not a Final GR Engine

The DMLG demonstrates that explicit geometry and modular recursion are necessary, but it is unlikely to be the final form of a GR engine suitable for AGI. Its canonical graph is powerful for narrative structure but limited in expressiveness, adaptability, and domain generality.

A full GR engine may require:

- multi-level canonical forms,
- dynamic paging and re-paging,
- structural compression and expansion,
- hierarchical recursion,
- and cross-domain geometric alignment.

Open question: *What is the minimal GR architecture that can scale from domain-bound intelligence to general intelligence while remaining interpretable and governable?*

8.4 Limitation 4: Recursion Depth Is Bounded by External Scaffolding

The pipeline performs recursion through retries, transitions, scoring, and keyword evolution, but the recursion is orchestrated externally by the narrator controller. The GR substrate itself does not yet perform deep internal recursion, self-evaluation, or multi-step reasoning.

Human recursive cognition is intrinsic, not scaffolded: reasoning loops, self-correction, and planning occur within the cognitive substrate itself.

Open question: *How can recursive depth be internalized within the GR module so that multi-step reasoning emerges natively rather than through external orchestration?*

8.5 Limitation 5: Domain-Bound Intelligence Is Not Yet Compositional

The narrator pipeline is a domain-bound AGI module: structurally complete but semantically restricted. It does not yet integrate multiple domain-bound modules into a unified compositional mind, nor does it coordinate heterogeneous domains through a mediator.

Open question: *How can multiple domain-bound AGI modules be composed into a coherent multi-domain cognitive architecture without losing the stability and identity of each module?*

8.6 Limitation 6: GR Scaling Remains an Open Frontier

The pipeline suggests that LLMs already provide sufficient informational (I) and energetic (E) capacity for AGI, and that the true bottleneck lies in engineered geometric and recursive structure (G and R). However, it remains unclear how far GR scaling can go, how it should be measured, and what structural thresholds must be crossed for general intelligence to emerge.

Open question: *What does it mean to “scale” G and R? What are the structural scaling laws for geometric and recursive intelligence, and how do they relate to the emergence of general-purpose cognition?*

8.7 Limitation 7: The Nature of Hybrid Self-Improvement Is Unresolved

The pipeline demonstrates weak hybrid self-improvement through keyword evolution, context accumulation, and synthetic distillation. But a full AGI requires a principled mechanism for structural self-improvement that is safe, interpretable, and bounded.

Open question: *How can a hybrid IEGR system improve its own GR substrate without uncontrolled drift, catastrophic forgetting, or loss of interpretability?*

8.8 Summary

The narrator pipeline shows that hybrid IEGR-complete systems can exhibit pre-AGI behaviour, but it also reveals the central challenge ahead: the design of a dynamic, self-improving GR engine with deep feedback loops and structural plasticity. LLMs already provide the IE roles of AGI; the frontier now lies in engineering the geometric and recursive substrate that can match the flexibility and self-organizing power of human cognition.

9 Why LLMs Feel Magical: IEGR Complementarity Between Humans and Machines

A final observation emerging from this work concerns the widespread intuition that large language models possess a kind of “magic”—a surprising breadth, fluency, and apparent intelligence that seems disproportionate to their mechanistic nature. The narrator pipeline provides a structural explanation for this phenomenon, grounded in the IEGR taxonomy.

Humans and LLMs occupy almost orthogonal regions of the IEGR space. Humans are strong in geometric structure (G) and recursive depth (R): they maintain stable world models, canonical conceptual structures, causal reasoning, long-horizon planning, and deep recursive self-correction. But humans are limited in informational richness (I) and energetic throughput (E): their knowledge is domain-bound, sensorily grounded, and energetically constrained.

LLMs exhibit the opposite profile. They are extraordinarily strong in I and E: they possess vast informational breadth, high-bandwidth semantic interpolation, and enormous energetic capacity for combinatorial inference. Yet they are weak in G and R: they lack stable geometry, canonical structure, persistent self-models, and intrinsic recursion.

The result is a profound structural complementarity:

Humans are GR-dominant systems with bounded IE. LLMs are IE-dominant systems with bounded GR. Together, they form an IEGR-complete hybrid intelligence.

This complementarity explains why LLMs feel powerful to humans: they supply the two cognitive roles in which humans are weakest. Conversely, humans supply the two roles in which LLMs are weakest. The narrator pipeline demonstrates that this complementarity is not merely philosophical but operational: when a structural generator (DMLG) provides G and R, and an

LLM provides I and E, the resulting system behaves like a miniature AGI—capable of planning, recursive refinement, world absorption, long-horizon coherence, and autonomous continuation.

This insight also reconciles the major positions in the AGI debate. Advocates of LLM scaling were correct that informational and energetic magnitude are necessary. Advocates of structural and symbolic approaches were correct that geometry and recursion are indispensable. And hybrid theorists were correct that general intelligence emerges only when both are present simultaneously.

The narrator pipeline therefore suggests a unifying conclusion:

AGI will not arise from scaling LLMs alone, nor from purely structural systems. It will arise from hybrid architectures in which LLMs supply the IE roles and engineered modules supply the GR roles.

In this sense, the system presented in this paper is not merely a narrative generator but a concrete demonstration of IEGR complementarity in action—a domain-bound AGI module whose behaviour illuminates the structural path toward general intelligence.

10 Conclusion

This paper introduced diffusion-driven narration as a generative methodology in which narrative worlds emerge through staged denoising across heterogeneous cognitive modules. While the immediate contribution lies in demonstrating a practical pipeline for autonomous long-form generation, the deeper significance of this work concerns what it reveals about the structure of intelligence itself.

The narrator pipeline shows that large language models already provide the informational richness (I) and energetic capacity (E) required for many aspects of intelligent behaviour. They supply semantic breadth, stylistic fluency, high-bandwidth inference, and evaluative judgment. Yet they lack the geometric structure (G) and recursive depth (R) necessary for stable world models, long-horizon planning, and autonomous continuation. These roles are supplied by the DMLG substrate, which provides canonical structure, modular paging, beam-search planning, and recursive governance.

Together, these components form an IEGR-complete hybrid system: a structural generator coupled to an informational evaluator. The resulting architecture exhibits pre-AGI properties—world absorption, recursive refinement, narrative planning, multi-agent blending, and autonomous continuation—despite relying only on existing models and modest computational resources. This suggests that the path toward AGI does not lie in scaling monolithic models indefinitely, but in engineering the geometric and recursive substrate that complements the strengths of LLMs.

This insight also clarifies a long-standing tension in the field. Advocates of LLM scaling were correct that informational and energetic magnitude are necessary. Advocates of structured and symbolic approaches were correct that explicit geometry and recursion are indispensable. And hybrid theorists were correct that general intelligence emerges only when both are present simultaneously. The narrator pipeline demonstrates that these perspectives are not contradictory but complementary: each captures a different dimension of the IEGR architecture that underlies intelligent systems.

At the same time, the work reveals the limitations of current approaches. The GR substrate remains largely static; its feedback loops are shallow; its structural learning is limited; and its recursion is externally orchestrated rather than intrinsic. A full AGI will require a dynamic, self-modifying, self-correcting GR engine capable of deep feedback integration and lifelong structural adaptation. The present system points toward this direction but does not yet reach it.

In this sense, diffusion-driven narration is both a demonstration and a beginning. It shows that hybrid IEGR-complete systems can already perform non-trivial autonomous cognition

within a bounded domain. It also identifies the frontier: the design of a dynamic geometric-recursive substrate that can scale beyond narrative generation toward general-purpose reasoning. The direction is clear, even if the destination has not yet been reached.

The work therefore contributes not only a generative methodology but a structural hypothesis about the nature of AGI: that intelligence arises from the recursive interaction of complementary cognitive roles, and that hybrid architectures—rather than monolithic models—provide the most promising path toward systems that are both powerful and governable. Diffusion-driven narration is one such system: a small but concrete step toward modular, interpretable, and structurally grounded artificial general intelligence.

A Triadic Cosmos Ecosystem Context

The present paper is part of the broader *Triadic Cosmos* ecosystem: a unified research program that develops a compositional ontology of intelligence together with modular, inspectable computational architectures. The ecosystem spans conceptual theory, executable reference implementations, and fully documented empirical demonstrations, forming a coherent framework for studying hybrid intelligent systems.

All materials, extended notes, and evolving documentation are publicly available in the open GitHub project: <https://github.com/triadic-cosmos>. Within the **triadic-system** repository, a complete reference implementation of the narrator pipeline—covering the distillation loop, multi-agent orchestration, and staged narrative diffusion—is provided under the AGPL-3.0 license. Complementary datasets, generated books, and trained micro-models are hosted in the separate **triadic-data** repository, enabling full reproducibility of the architecture and its outputs.

This paper does not present implementation details or narrative samples directly; instead, it situates the method within the conceptual and architectural framework. Concrete demonstrations of the system in action are developed in two companion volumes within the ecosystem:

Tales of a Sophisticated Parrot

A microscopic, fully transparent demonstration of hybrid intelligence. This volume documents a structurally complete IEGR system built from a compact 4-bit LLM and a set of tiny DMLG world models. It presents seven complete, unredacted books generated by the system—canonical artefacts produced through deterministic structure interacting with expressive refinement. The book serves as an empirical atlas of *sophisticated parrots*: hybrid agents that combine energetic and informational roles (LLM) with geometric and recursive roles (DMLG), illustrating how coherent long-form literature can emerge at minimal scale.

More Tales of a Sophisticated Parrot: Diffusion-Driven Narration

A comprehensive exploration of diffusion as a narrative process. This volume introduces the generative methodology formalized in the present paper: narrative structure emerging through staged denoising across heterogeneous cognitive modules. It demonstrates six generative modalities—Sampling, Remaster, Blending, Blending + Remaster, Narration, and Narration + Director’s Cut—and shows how each corresponds to a distinct IEGR transformation within the diffusion trajectory. The book expands the methodology into a general framework for *synthetic narratology*, where worlds are not written but grown through structured transformations of latent narrative geometry.

Together, the repositories and the two companion books form the empirical substrate of the *Triadic Cosmos*: a research ecosystem where conceptual theory, executable architecture, and narrative artefacts reinforce one another. The present paper provides the methodological and taxonomic foundation for this ecosystem, formalizing *Diffusion-Driven Narration* as a modular, hybrid-AGI approach to autonomous long-form generation.

References

- Ammanabrolu, Prithviraj and Mark Riedl (2020). “Story Realization: Expanding Plot Events into Sentences”. In: *AAAI*.
- Besold, Tarek et al. (2017). “Neural-Symbolic Learning and Reasoning: A Survey and Interpretation”. In: *arXiv preprint arXiv:1711.03902*.
- Brown, Tom, Benjamin Mann, Nick Ryder, et al. (2020). “Language Models are Few-Shot Learners”. In: *NeurIPS*.
- Chen, Angelica et al. (2023). “RE3: Reasoning via Rewriting”. In: *ICLR*.
- De Witte, Joachim (2026a). “A Dynamic and Modular Architecture for Language Modeling”. In: *Triadic Cosmos Technical Report*. Available at: <https://github.com/triadic-cosmos>.
- (2026b). “A Structural Taxonomy of Intelligence: IEGR-Scale, Parrot Regimes, and the Conditions for Human-Like AGI”. In: *Triadic Cosmos Technical Report*. Available at: <https://github.com/triadic-cosmos>.
- (2026c). “Language as Graph Traversal: The Dynamic Modular Language Graph (DMLG)”. In: *Triadic Cosmos Technical Report*. Available at: <https://github.com/triadic-cosmos>.
- Ganguli, Deep, Amanda Askell, Yuntao Bai, et al. (2022). “Red Teaming Language Models to Reduce Harms”. In: *arXiv preprint arXiv:2209.07858*.
- Gong, Yihan et al. (2023). “DiffuSeq: Sequence to Sequence Text Generation with Diffusion Models”. In: *ICLR*.
- Hokamp, Chris and Qun Liu (2017). “Lexically Constrained Decoding for Sequence Generation”. In: *ACL*.
- Jurafsky, Daniel and James Martin (2023). *Speech and Language Processing*. Prentice Hall.
- Li, Xiang et al. (2022). “Diffusion-LM Improves Controllable Text Generation”. In: *NeurIPS*.
- Madaan, Aman, Amir Yazdanbakhsh, et al. (2023). “Self-Refine: Iterative Refinement with Self-Feedback”. In: *arXiv preprint arXiv:2303.17651*.
- Touvron, Hugo, Thibaut Lavril, Gautier Izacard, et al. (2023). “LLaMA 2: Open Foundation and Fine-Tuned Chat Models”. In: *arXiv preprint arXiv:2307.09288*.
- Wei, Jason, Xuezhi Wang, Dale Schuurmans, et al. (2022). “Chain-of-Thought Prompting Elicits Reasoning in Large Language Models”. In: *NeurIPS*.
- Wu, Yujia et al. (2023). “AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation”. In: *arXiv preprint arXiv:2308.08155*.
- Zeng, Andy, Pete Florence, et al. (2022). “Socratic Models: Composing Zero-Shot Multimodal Reasoning with Language”. In: *ICLR*.
- Zhou, Wangchunshu et al. (2023). “A Survey of Hierarchical Text Generation”. In: *arXiv preprint arXiv:2302.09170*.