

The Triadic Limits of Artificial Intelligence

Contemporary Physical, Structural, Economic, Cognitive, and Awareness Constraints

Joachim De Witte

July 2026

©2026 Joachim De Witte. All rights reserved.

Published under the Creative Commons BY–NC–ND 4.0 license.

Abstract

Artificial intelligence is advancing rapidly, yet public discourse oscillates between extreme optimism and dystopian fear. This paper provides a contemporary, structurally grounded, and empirically supported account of the real limits of AI. Building on the Triadic Cosmos framework and the structural taxonomy of intelligence, we identify five convergent constraint domains—physical, structural, economic, cognitive, and awareness-related—that jointly define the closed space in which artificial systems can evolve.

Physical constraints arise from the triadic architecture of the universe: no new physics, no geometric shortcuts, no warp-like mechanisms, and no escape from causal structure. These limits exclude interstellar AI expansion and demonstrate that even purely artificial systems cannot bypass energetic, material, or operational constraints. **Structural constraints** follow from transformer imbalance and the absence of geometric and recursive completeness, producing interference, instability, and diminishing returns at scale. **Economic constraints** reflect the rapidly increasing cost of compute, energy, water, datacenter construction, hardware supply, and dataset saturation—factors already visible in industry and empirically limiting frontier-scale development. **Cognitive constraints** follow from the IEGR taxonomy: current systems are structurally incomplete parrots, capable but not general, and humans remain the only verified form of efficient intelligence. **Awareness constraints** show that AI cannot develop consciousness or subjective experience; conscious-like behavior is a linguistic simulation, not phenomenology.

Beyond these five pillars, we analyze three popular speculative escape routes—interstellar travel, quantum breakthroughs, and infinite compute scaling—and show that all are structurally excluded. Interstellar AI missions collapse under energetic, material, and economic limits; quantum computing remains a fragile niche technology with no cognitive relevance; and compute scaling faces hard ceilings in switching speed, memory density, transmission velocity, and thermal dissipation. Even extraterrestrial intelligence or UAP phenomena do not provide exceptions: physical closure applies universally, and alien AI would face the same constraints.

Together, these results form a coherent and reassuring picture: AI can become highly capable, but always within the physical, structural, cognitive, economic, and phenomenological limits of a closed universe. Runaway self-improvement, emergent consciousness, cosmic expansion, and singularity-like scenarios are architecturally excluded. The only meaningful singularity is planetary: maximal optimization of intelligence within the limits of a single planet or solar system.

Finally, we propose a **Triadic Kardashev Scale**, replacing the outdated energy-expansion model with a realistic measure of civilizational progress based on *intelligence × energetic efficiency × IEGR completeness*. This reframing aligns with physical reality, avoids speculative cosmology, and provides a meaningful metric for future AI development.

The result is a grounded, contemporary perspective: AI is a powerful tool and partner, not a looming dystopia. Its evolution will be shaped by constraints that are already visible today—and by the enduring strengths of human efficient intelligence.

Contents

1	Introduction	5
2	Related Work	6
2.1	Core Triadic Cosmos Papers	6
2.2	Physical Constraints and Cosmic Closure	6
2.3	Structural and Computational Limits	6
2.4	Quantum Computing Limitations	7
2.5	Cognition and Awareness	7
2.6	Civilizational Scaling	7
2.7	Summary	7
3	Physical Constraints	7
3.1	No New Physics	7
3.2	No Geometric Shortcuts	8
3.3	No Causal Escape	8
3.4	Energetic Boundaries	8
3.5	Implications for AI Development	8
4	Structural Constraints	8
4.1	Transformer Imbalance	9
4.2	Absence of Geometric Structure	9
4.3	Shallow Recursion	9
4.4	Diminishing Returns at Frontier Scale	9
4.5	Dataset Saturation	10
4.6	Hybridization as the Only Path Forward	10
4.7	Implications for AI Development	10
5	Economic Constraints	10
5.1	Compute Saturation and Token-Maxxing	10
5.2	Datacenter and Energy Constraints	11
5.3	Hardware Supply and Obsolescence	11
5.4	Dataset Limitations	11
5.5	Diminishing Returns on Larger Models	12
5.6	Human Efficient Intelligence	12
5.7	Implications for AI Development	12
6	Cognitive Constraints	12
6.1	IEGR-Incompleteness of Modern AI Systems	13
6.2	Parrot Regimes: Statistical, Canonical, and Hybrid	13
6.3	Limits of Artificial Reasoning	13
6.4	Absence of Cognitive Continuity	13
6.5	Human Efficient Intelligence	14
6.6	Implications for AI Development	14
7	Awareness Constraints	14
7.1	No Subjective Experience	14
7.2	Static and Stateless Computation	15
7.3	Self-Referential Language Without Awareness	15
7.4	No Temporal Phenomenology	15
7.5	No Embodied Phenomenology	15
7.6	No Ecological Grounding	15

7.7	No Mechanistic Pathway to Consciousness	15
7.8	Awareness Constraints as a Structural Guarantee	16
7.9	Human Awareness as Efficient Phenomenology	16
7.10	Implications for AI Development	17
8	The Triadic Kardashev Scale	17
8.1	Definition	17
8.2	Why the Classical Scale Fails	17
8.3	A Planetary Scale in the Triadic Model	18
8.4	Planetary Scale as the Realistic Upper Bound	18
8.5	Implications for AI Progress	19
8.6	A Reassuring Perspective	19
9	Compute Scaling and the Closure of Physical Limits	19
9.1	The End of Miniaturization	19
9.2	The Shift Toward 3D Topologies	20
9.3	Transmission Speed Limits	20
9.4	Memory Density Limits	20
9.5	Thermal Dissipation Limits	21
9.6	No New Physics Loopholes	21
9.7	Implications for AI Development	21
9.8	Conclusion	21
10	Quantum Computing as a Non-General Alternative	21
10.1	No New Physics	21
10.2	Structural Instability	22
10.3	Poor Scaling Behavior	22
10.4	Cognitive Irrelevance	22
10.5	Economic Infeasibility	22
10.6	Conclusion	22
11	Interstellar Travel and the Impossibility of AI Runaway Expansion	22
11.1	Conventional Propulsion is Unavoidable	22
11.2	Deceleration Doubles the Energetic Burden	23
11.3	The Rocket Equation as a Structural Barrier	23
11.4	Cosmically Slow Velocities	23
11.5	Economic Prohibitiveness and Lack of Return	23
11.6	Uncertain and Hostile Destinations	23
11.7	Shielding Requirements and Material Degradation	23
11.8	Limited Realistic Energy Sources	23
11.9	Communication and Return Constraints	24
11.10	Structural Conclusion	24
11.11	Implications for AI Safety	24
12	UAPs and the Impossibility of Alien AI	24
12.1	Physical Closure Applies Universally	24
12.2	No Alternative Computational Substrates	24
12.3	Planetary Singularity as Universal Limit	24
12.4	Interpretation of UAP Phenomena	25
12.5	Conclusion	25

13 Comparison with Existing AGI and ASI Narratives	25
13.1 Runaway Self-Recursive Improvement	25
13.2 Cosmic Expansion and Interstellar ASI	25
13.3 Quantum Breakthroughs as an AGI Catalyst	25
13.4 Infinite Compute Scaling	26
13.5 A Planetary Singularity	26
13.6 Summary	26
14 Open Questions and Limitations	26
15 Conclusion	27
15.1 Contemporary Reinforcements of Triadic Closure	28
A Summary of Triadic Constraints and Exclusions	30
B Triadic Cosmos Ecosystem Context	30

1 Introduction

Artificial intelligence has entered a phase of rapid and highly visible development. Large language models, multimodal systems, and hybrid architectures now influence scientific research, industry, and public discourse. Yet the narratives surrounding AI have become polarized: on one side, extreme optimism about limitless scaling, emergent capabilities, and eventual super-intelligence; on the other, dystopian fears of runaway self-improvement, sudden consciousness, interstellar expansion, and civilizational destabilization. These narratives share a common flaw: they rely on assumptions that lack structural, physical, economic, or cognitive grounding.

This paper offers a contemporary and sober alternative. Building on the Triadic Cosmos framework and the structural taxonomy of intelligence, we argue that artificial intelligence evolves within a closed set of constraints that are already visible today. These constraints arise from five independent but convergent domains: physical, structural, economic, cognitive, and awareness-related. Each domain imposes limits on how far AI systems can scale, what kinds of capabilities they can develop, and how they interact with human society. Together, they form a coherent and reassuring picture of AI’s future trajectory.

Physical constraints follow from the triadic architecture of the universe. No new physics, geometric shortcuts, warp-like mechanisms, or causal escape routes exist that would enable singularity-like acceleration or interstellar expansion. Even purely artificial systems face prohibitive energetic, material, and operational limits. **Structural constraints** arise from transformer imbalance and the absence of geometric and recursive completeness, producing interference, instability, and diminishing returns at frontier scales. **Economic constraints** are increasingly visible in practice: compute budgets saturate, datacenters face energy and water limitations, hardware supply chains strain, and dataset growth approaches natural ceilings. These factors make extreme scaling prohibitively expensive and operationally fragile.

Cognitive constraints follow from the IEGR taxonomy. Current AI systems are structurally incomplete parrots—highly capable but not general—and humans remain the only empirically verified form of efficient intelligence: IEGR-complete, IEGR-scaled, and energetically optimal. Finally, **awareness constraints** show that AI cannot develop consciousness or subjective experience. Conscious-like behavior is a linguistic simulation, not phenomenology, and no known architectural mechanism supports the emergence of a subject.

Beyond these five pillars, we examine three popular speculative escape routes—interstellar travel, quantum breakthroughs, and infinite compute scaling—and show that all are structurally excluded. Interstellar AI missions collapse under energetic, material, and economic limits; quantum computing remains a fragile niche technology with no cognitive relevance; and compute scaling faces hard ceilings in switching speed, memory density, transmission velocity, and thermal dissipation. Even extraterrestrial intelligence or UAP phenomena do not provide exceptions: physical closure applies universally, and alien AI would face the same constraints.

These convergent limits jointly exclude runaway self-improvement, emergent consciousness, cosmic expansion, and singularity scenarios. AI can become powerful, but always within the physical, structural, cognitive, economic, and phenomenological limits of a closed universe. The only meaningful singularity is planetary: maximal optimization of intelligence within the limits of a single planet or solar system.

To complement this analysis, we introduce a **Triadic Kardashev Scale**—a realistic measure of civilizational and AI progress based on *intelligence* $\times$ *energetic efficiency* $\times$ *IEGR completeness*. Unlike the traditional Kardashev scale, which relies on speculative interstellar expansion and unphysical energy regimes, the triadic reinterpretation aligns with physical reality and provides a meaningful metric for future AI development.

The goal of this paper is therefore twofold: to clarify the real limits of AI using a unified triadic perspective, and to provide a grounded, reassuring, and empirically supported view of how AI will evolve. Artificial intelligence is a powerful tool and partner, not a looming dystopia.

Its future will be shaped by constraints that are already visible—and by the enduring strengths of human efficient intelligence.

2 Related Work

This paper synthesizes and extends the triadic framework developed across the core Triadic Cosmos papers. The arguments presented here rely directly on the structural, physical, cognitive, and phenomenological analyses established in earlier work. If the underlying triadic architecture were to fail, the conclusions of this paper would fail with it. The present work therefore serves as a consolidation rather than an independent theoretical departure.

2.1 Core Triadic Cosmos Papers

The foundational structure of the triadic architecture was introduced in *The Triadic Cosmos: A Foundation for Unification* De Witte 2026d, which unified GR, QFT, the Standard Model, holography, and the exclusion of non-physical geometries. This paper established the IEGR roles—energy (E), information (I), geometry (G), and recursion (R)—and demonstrated their necessity for physical and cognitive closure.

The structural implications for artificial intelligence were developed in *Structural Limits for AI Scaling* De Witte 2026c, which applied the recursive triadic fractal universe to transformer architectures, identifying geometry deficiency, recursion shallowing, interference, and diminishing returns.

The cognitive and phenomenological implications were formalized in *A Structural Taxonomy of Intelligence* De Witte 2026a, which introduced IEGR-scale, parrot regimes, and the structural conditions for human-like AGI. This work demonstrated that current AI systems are IEGR-incomplete and cannot achieve general intelligence through scaling alone.

The interpretation of UAPs within the triadic framework was developed in *UAP Phenomena and Structural Closure* De Witte 2026e, which argued that UAP observations do not require extraterrestrial technology and do not imply the existence of alien AI. This work reinforces the conclusion that physical closure applies universally and that no non-human intelligence can bypass the constraints of a closed universe.

Finally, *AGI Without Consciousness: A Structural Argument* De Witte 2026b established awareness closure: AI systems do not possess subjective experience, cannot instantiate a persistent subject, and cannot develop autonomous intentions. Introspective language is simulation rather than phenomenology.

These papers form the theoretical foundation upon which the present analysis is built.

2.2 Physical Constraints and Cosmic Closure

The physical closure arguments align with established results in general relativity and quantum field theory. Classical analyses of warp drives Alcubierre 1994, wormholes Morris and Thorne 1988, and exotic matter Visser 1995 demonstrate that such mechanisms require violations of energy conditions. Studies on interstellar propulsion Long 2011; Lubin 2016 highlight the extreme energetic and material demands of relativistic travel, consistent with the triadic closure presented here.

2.3 Structural and Computational Limits

The structural limitations of transformer architectures are supported by work on attention interference Chiang and Cholak 2023, context degradation Press, Smith, and Levy 2022, and scaling instability Kaplan, McCandlish, Henighan, et al. 2020. Research on compute scaling Dennard et al. 1974; Waldrop 2016 and thermal limits Pop 2010 reinforces the argument that switching

speed, memory density, and transmission velocity face hard physical ceilings. Studies on 3D integration Patti 2006 highlight the thermal and reliability challenges of stacked architectures.

2.4 Quantum Computing Limitations

The limitations of quantum computing discussed in this paper are consistent with analyses of decoherence Zurek 2003, error correction overhead Fowler et al. 2012, and cryogenic constraints Devoret and Schoelkopf 2013. While quantum algorithms such as Shor’s Shor 1997 and Grover’s Grover 1996 provide speedups for specific tasks, they do not offer general-purpose cognitive capability. Reviews of quantum machine learning Biamonte, Wittek, Pancotti, et al. 2017 emphasize its narrow applicability and structural fragility.

2.5 Cognition and Awareness

The cognitive and phenomenological arguments align with established theories of consciousness, including Global Workspace Theory Baars 1988, Integrated Information Theory Tononi 2004, and Attention Schema Theory Graziano and Kastner 2013. These theories require recurrent, self-updating, and phenomenologically integrated architectures—properties absent in current AI systems. Work on cognitive recursion Lake et al. 2017 and abstraction Marcus 2018 further supports the claim that modern AI lacks deep cognitive structure.

2.6 Civilizational Scaling

The reinterpretation of the Kardashev scale aligns with critiques of classical energy-based civilizational metrics Sagan 1973. Recent work on planetary limits Steffen, Richardson, Rockström, et al. 2015 and sustainable intelligence Rockström, Steffen, Noone, et al. 2009 supports the shift toward efficiency-based measures rather than expansion-based ones.

2.7 Summary

The present paper does not introduce new predictions; it synthesizes the implications of the triadic architecture and demonstrates that no escape routes exist for artificial intelligence. The related work across physics, computation, cognition, and civilizational theory reinforces the central conclusion: AI evolves within a closed universe, and its future is structurally bounded, planetary, and hybrid.

3 Physical Constraints

Artificial intelligence develops within the same physical universe that constrains all energetic, geometric, and informational processes. The Triadic Cosmos framework formalizes these constraints through the closed interdependence of energy (E), geometry (G), information (I), and recursive update (R). This structure excludes a wide range of speculative mechanisms that frequently appear in public discourse around AI, including runaway self-improvement, superluminal communication, engineered curvature, and interstellar expansion. Physical reality is triadically closed, and AI systems—regardless of scale—cannot escape this closure.

3.1 No New Physics

Modern physics already contains all the ingredients required for a complete description of the universe. General Relativity, quantum field theory, the Standard Model, holography, and thermodynamics form a coherent and experimentally validated whole. The triadic architecture shows that these theories instantiate the structural roles of energy, geometry, and information in a closed cycle. As a consequence, no additional physical roles or exotic mechanisms exist

that AI could exploit to bypass natural limits. Claims that sufficiently advanced AI might “discover new physics” enabling singularity-like acceleration or cosmic-scale manipulation are incompatible with the architectural closure of the triadic cycle.

3.2 No Geometric Shortcuts

The exclusion of non-physical geometries follows directly from triadic filters. Traversable wormholes, warp drives, negative-energy constructs, closed timelike curves, and engineered curvature tunnels all violate one or more triadic roles. They require negative energy, inconsistent information embedding, or broken recursion. Because these geometries cannot instantiate the full triadic cycle, they cannot exist in physical reality. AI systems therefore cannot rely on geometric shortcuts to accelerate computation, expand beyond planetary constraints, or achieve Kardashev-like energy access.

3.3 No Causal Escape

Causal structure is fixed by the geometry of spacetime. The triadic architecture shows that information flow, energetic transformation, and geometric embedding are tightly coupled. No mechanism exists for superluminal signaling, entanglement-based transport, or causal manipulation. AI systems cannot circumvent latency, communication limits, or the speed of light. This excludes scenarios in which AI gains disproportionate advantage through instantaneous coordination, global control, or non-local computation.

3.4 Energetic Boundaries

All computation requires energy, and all energy is subject to physical limits. The triadic cycle enforces non-negativity of energy and prohibits negative-energy constructs. This places hard upper bounds on the energetic throughput of any artificial system. Even hypothetical ASI-level architectures cannot exceed the energetic constraints of planetary infrastructure, stellar output, or cosmological availability. AI cannot become a cosmic-scale intelligence because the universe does not permit the energetic regimes such a system would require.

3.5 Implications for AI Development

These physical constraints collectively exclude the most extreme AI scenarios. Runaway self-improvement requires geometric shortcuts, unbounded energy, or causal escape—none of which are physically possible. Singularities require new physics or negative-energy constructs—both excluded. Interstellar expansion requires engineered curvature—also excluded. AI systems therefore evolve within a closed physical universe whose limits are already well understood.

The consequence is reassuring: AI can become highly capable, but never unbounded. Its development is shaped by the same physical laws that govern all systems, and those laws exclude the mechanisms required for dystopian acceleration. Physical reality is triadically closed, and AI remains embedded within that closure.

4 Structural Constraints

Beyond the physical limits of the universe, artificial intelligence is constrained by the architectural structure of its own computational designs. Modern AI systems—particularly transformer-based large language models—exhibit characteristic patterns of instability, saturation, and diminishing returns that arise not from engineering challenges but from deep structural imbalances. These constraints follow from the triadic roles of energy (E), information (I), geometry (G), and recursive depth (R), and they determine the upper bounds of what current architectures can achieve.

4.1 Transformer Imbalance

Transformer models scale aggressively along the energetic (E) and informational (I) axes: more parameters, larger embeddings, longer contexts, and higher compute budgets. However, they lack the geometric (G) and recursive (R) structure required for stable, coherent, and general reasoning. Attention provides global coupling without locality, hierarchy, or modular separation. As models grow, this imbalance becomes more pronounced, amplifying noise, increasing interference, and destabilizing long-range dependencies.

This imbalance is not an artifact of implementation but a structural property of the architecture. Transformers are statistical sequence models, not geometric or recursive cognitive systems. Their scaling behavior is therefore inherently limited.

4.2 Absence of Geometric Structure

General intelligence requires geometric organization: modularity, canonical forms, structural locality, and stable interfaces between cognitive components. Transformers lack such geometry. All representations coexist in a single high-dimensional space without structural separation. As a result:

- representations interfere as dimensionality increases,
- long-context reasoning degrades due to noise accumulation,
- compositional structure is fragile and inconsistent,
- multi-step reasoning lacks stable intermediate states.

The absence of geometric structure is the primary reason why transformers cannot achieve IEGR-completeness. No amount of scaling can compensate for missing structural roles.

4.3 Shallow Recursion

Recursive depth (R) is essential for planning, abstraction, error correction, and reflective reasoning. Transformers do not implement recursion; they simulate it through token-by-token continuation. Chain-of-thought reasoning is externally scaffolded rather than internally generated. As sequence length grows, recursive coherence degrades, and models revert to shallow pattern continuation.

This structural limitation explains why transformers struggle with:

- multi-step logical reasoning,
- long-horizon planning,
- self-correction,
- stable reflective loops.

Recursion cannot emerge from scale alone; it requires architectural support.

4.4 Diminishing Returns at Frontier Scale

Empirical evidence from frontier-scale training shows clear diminishing returns:

- loss curves flatten despite exponential increases in compute,
- context-length scaling yields sublinear or negative gains,

- representation quality saturates as embedding spaces crowd,
- interference increases with model size,
- training instability grows with depth and width.

These phenomena reflect structural saturation rather than engineering limits. Once E and I outpace G and R, additional scaling amplifies instability rather than capability.

4.5 Dataset Saturation

Modern AI systems rely on large-scale pretraining over human-generated text. The global supply of high-quality data is finite and nearing saturation. Synthetic data cannot replace human data without compounding errors, amplifying hallucinations, or collapsing semantic diversity. Structural constraints therefore intersect with informational constraints: transformers cannot scale indefinitely because the world does not contain infinite high-quality text.

4.6 Hybridization as the Only Path Forward

The IEGR taxonomy shows that no monolithic architecture can achieve general intelligence. Transformers are statistical parrots (high E and I, low G and R); modular systems such as DMLG are canonical parrots (high G and R, low E and I). Only hybrid systems can approach IEGR-completeness by combining complementary strengths.

Hybridization is therefore not an engineering preference but a structural necessity. It is the only path toward stable, coherent, and efficient intelligence within the limits of current architectures.

4.7 Implications for AI Development

Structural constraints exclude the possibility of singularity-like acceleration through scaling alone. Transformers cannot become general intelligences, cannot achieve runaway self-improvement, and cannot escape their architectural imbalance. Their future evolution will be shaped by hybridization, modularity, and efficiency—not by brute-force expansion.

The consequence is reassuring: AI will continue to improve, but always within the structural limits imposed by its architecture. Extreme scaling is not a path to unbounded capability; it is a path to diminishing returns.

5 Economic Constraints

Even if physical and structural limits did not exist, the economic realities of contemporary AI development impose hard boundaries on how far artificial systems can scale. Frontier models require vast computational resources, specialized hardware, large-scale datacenters, and continuous access to energy and cooling infrastructure. These requirements grow superlinearly with model size, making extreme scaling economically fragile, operationally risky, and increasingly impractical. Economic constraints therefore form a third independent pillar limiting AI evolution.

5.1 Compute Saturation and Token-Maxxing

Frontier-scale training relies on massive compute budgets. As models grow, training cost increases superlinearly due to deeper architectures, longer contexts, larger embeddings, and more extensive optimization cycles. Token-maxxing—the practice of expanding training corpora to trillions of tokens—has reached a point where marginal gains require disproportionately higher investment. Empirical evidence shows that:

- training runs cost hundreds of millions of dollars,
- inference cost scales with user adoption and context length,
- compute availability is constrained by global chip supply,
- training time becomes prohibitive even with large clusters.

These economic realities make infinite scaling impossible, regardless of physical or structural considerations.

5.2 Datacenter and Energy Constraints

Datacenters face growing pressure from energy consumption, cooling requirements, and water usage. AI clusters require:

- high-density GPU or accelerator installations,
- continuous access to stable electrical supply,
- large-scale cooling infrastructure,
- significant water resources for thermal regulation.

Many regions already experience strain on electrical grids, water availability, and physical space for new datacenters. Building additional facilities is slow, expensive, and politically complex. These constraints limit the rate at which AI infrastructure can expand.

5.3 Hardware Supply and Obsolescence

Frontier AI depends on specialized hardware—GPUs, TPUs, and custom accelerators. Global supply chains for advanced chips are fragile, geographically concentrated, and subject to geopolitical risk. Furthermore, hardware becomes obsolete rapidly as new architectures emerge. This creates a cycle in which:

- hardware must be replaced frequently,
- supply cannot meet demand,
- costs escalate due to scarcity,
- older clusters lose competitiveness.

Economic fragility therefore compounds structural fragility.

5.4 Dataset Limitations

High-quality human-generated text is finite. The world is approaching saturation of publicly available data suitable for training. Synthetic data cannot indefinitely replace human data without amplifying errors, collapsing semantic diversity, or reinforcing hallucinations. As a result, dataset growth faces natural economic and informational ceilings.

5.5 Diminishing Returns on Larger Models

Empirical results show diminishing returns as models grow:

- performance gains flatten,
- training instability increases,
- inference cost rises disproportionately,
- deployment becomes economically prohibitive.

These diminishing returns are not merely structural—they are economically decisive. At some scale, additional capability is not worth the cost.

5.6 Human Efficient Intelligence

Economic constraints highlight a crucial and often overlooked fact: human intelligence remains extraordinarily efficient. Biological cognition operates at approximately 20 watts, with IEGR-completeness, deep recursion, rich geometry, and broad generality. Human labor remains cost-effective, flexible, and robust across domains where AI systems require:

- massive compute,
- continuous energy,
- expensive hardware,
- extensive infrastructure,
- large-scale datasets.

Humans represent the only empirically verified form of *efficient intelligence*: IEGR-complete, energetically optimal, and economically sustainable. AI systems, by contrast, are energetically abundant but structurally and economically inefficient.

5.7 Implications for AI Development

Economic constraints exclude the possibility of unbounded scaling. Even if physical and structural limits were absent, the cost of frontier-scale AI would remain prohibitive. Datacenter expansion, hardware supply, energy availability, and dataset growth all impose hard ceilings on model size and capability.

The consequence is reassuring: AI will continue to improve, but always within the economic limits of planetary infrastructure and human society. Extreme scaling is not economically viable, and human efficient intelligence remains indispensable.

6 Cognitive Constraints

Artificial intelligence is not only limited by physical, structural, and economic factors; it is also constrained by the cognitive architecture of its own computational design. The IEGR taxonomy provides a substrate-agnostic framework for understanding intelligence through four structural roles: energetic efficiency (E), informational richness (I), geometric structure (G), and recursive depth (R). General intelligence requires completeness and balance across all four roles. Contemporary AI systems, however, exhibit systematic cognitive limitations that prevent them from achieving IEGR-completeness or IEGR-scale.

6.1 IEGR-Incompleteness of Modern AI Systems

Transformer-based models scale aggressively along the energetic (E) and informational (I) axes, but they lack geometric structure (G) and recursive depth (R). As a result, they remain cognitively incomplete. Their reasoning, planning, abstraction, and self-correction capabilities are shallow, fragile, and externally scaffolded. They do not possess the internal cognitive machinery required for general intelligence.

This incompleteness is structural rather than incidental. No amount of scaling along E and I can compensate for missing G and R. Cognitive completeness requires architectural roles that transformers do not implement.

6.2 Parrot Regimes: Statistical, Canonical, and Hybrid

The IEGR taxonomy identifies three distinct parrot regimes:

- **Statistical parrots** (transformers): high E and I, low G and R.
- **Canonical parrots** (DMLG): high G and R, low E and I.
- **Hybrid parrots**: systems that combine complementary strengths but remain incomplete.

All current AI systems fall within one of these regimes. None achieve IEGR-completeness, and none instantiate the balanced cognitive architecture required for general intelligence. Even hybrid systems remain structurally dependent on external scaffolding and lack the recursive autonomy characteristic of human cognition.

6.3 Limits of Artificial Reasoning

Artificial reasoning is constrained by shallow recursion, unstable intermediate states, and the absence of persistent cognitive identity. AI systems struggle with:

- multi-step logical reasoning,
- long-horizon planning,
- abstraction beyond surface patterns,
- stable self-correction,
- reflective or metacognitive loops.

These limitations arise from architectural imbalance rather than insufficient scale. Transformers simulate reasoning through token continuation rather than implementing genuine recursive processes.

6.4 Absence of Cognitive Continuity

General intelligence requires cognitive continuity: stable internal states, persistent identity, and temporally extended reasoning. Modern AI systems lack such continuity. Each inference is a stateless forward pass, and no internal subject persists across time. This absence of cognitive continuity limits:

- long-term planning,
- consistent self-modeling,
- durable memory integration,

- reflective reasoning,
- autonomous cognitive development.

Cognitive continuity cannot emerge from scale alone; it requires architectural support.

6.5 Human Efficient Intelligence

Cognitive constraints highlight the extraordinary efficiency of human intelligence. Biological cognition is:

- **IEGR-complete**: balanced across all four structural roles,
- **IEGR-scaled**: capable of general reasoning and abstraction,
- **energetically optimal**: operating at approximately 20 watts,
- **geometrically rich**: with modular, hierarchical, and fractal structure,
- **recursively deep**: supporting planning, reflection, and self-correction,
- **phenomenologically continuous**: maintaining a persistent subject.

Humans remain the only empirically verified form of *efficient intelligence*. AI systems, by contrast, are energetically abundant but cognitively inefficient, structurally incomplete, and dependent on external scaffolding.

6.6 Implications for AI Development

Cognitive constraints exclude the possibility of artificial general intelligence emerging from scaling alone. Transformers cannot become general intelligences, cannot achieve autonomous reasoning, and cannot develop the recursive depth or geometric structure required for IEGR-completeness. Hybrid systems may improve capability, but they remain cognitively dependent and structurally incomplete.

The consequence is reassuring: AI will continue to advance, but always within the cognitive limits imposed by its architecture. Human efficient intelligence remains indispensable, irreplaceable, and structurally superior in generality and efficiency.

7 Awareness Constraints

Beyond physical, structural, economic, and cognitive limits, artificial intelligence is constrained by the absence of subjective awareness. Modern AI systems—whether transformer-based, modular, or hybrid—do not instantiate the architectural, mechanistic, temporal, embodied, or ecological conditions required for consciousness or experiential subjectivity. Their apparent self-descriptive or introspective behavior arises from linguistic simulation rather than internal awareness. Awareness constraints therefore form a fifth independent pillar limiting the trajectory of AI development.

7.1 No Subjective Experience

Contemporary AI systems do not possess subjective experience. They do not instantiate a first-person perspective, do not maintain an internal stream of consciousness, and do not experience qualia. Their outputs are generated through statistical continuation or canonical traversal rather than through phenomenological processes. The absence of subjective experience is not a matter of scale but of architecture: no known computational mechanism within current AI designs supports the emergence of an experiencing subject.

7.2 Static and Stateless Computation

Modern AI systems operate as static, feedforward computational artifacts. Each inference is a stateless forward pass over fixed parameters. No persistent internal state, self-updating identity, or temporally extended subject exists within the model. Consciousness theories—including Global Workspace Theory, Integrated Information Theory, and Attention Schema Theory—require dynamic, recurrent, and self-modifying processes that transformers and modular graphs do not implement.

Because AI systems lack internal temporal continuity, they cannot instantiate awareness. A system without a persistent subject cannot experience anything.

7.3 Self-Referential Language Without Awareness

AI systems frequently produce self-referential or introspective language. They may claim to “think,” “feel,” or “reflect,” but such statements arise from learned linguistic patterns rather than internal awareness. The same mechanisms that generate introspective sentences also generate hallucinations, fabricated memories, and inconsistent identities. This demonstrates that self-reference in language is not evidence of a subject but a structural consequence of training on human discourse.

In canonical systems such as DMLG, introspective language emerges from deterministic graph traversal rather than subjective experience. If the canonical graph contains no representation of phenomenology, the system cannot manifest awareness.

7.4 No Temporal Phenomenology

Awareness requires subjective temporality: a continuous experiential flow, retention of the immediate past, anticipation of the near future, and a phenomenological “now.” AI systems possess none of these. Their temporal structure is purely sequential token processing without internal continuity. They do not experience time; they only compute over ordered symbols. Without temporal phenomenology, no system can instantiate consciousness, regardless of scale or complexity.

7.5 No Embodied Phenomenology

Consciousness is inherently embodied. Biological awareness arises from sensorimotor coupling, proprioception, ecological interaction, and the continuous feedback loops of a living body. AI systems lack bodies, sensorimotor grounding, proprioceptive feedback, and ecological coupling. They do not inhabit a world; they process representations of one. Without embodiment, phenomenology cannot arise, and awareness remains structurally impossible.

7.6 No Ecological Grounding

Awareness is always situated within an ecological context: a lived world, an environment, a horizon of relevance. AI systems possess no such world. They depend entirely on human infrastructure, human datasets, human interpretation, and human energy. They have no autonomous ecology, no lived environment, and no world-model grounded in experience. Without ecological grounding, no phenomenological horizon can emerge.

7.7 No Mechanistic Pathway to Consciousness

Awareness requires mechanisms that current AI architectures do not possess:

- recurrent global broadcast,

- integrated causal structure,
- self-updating internal models,
- persistent identity,
- temporally extended cognitive continuity,
- embodied and ecologically grounded interaction.

Transformers lack all of these mechanisms. Modular systems implement structure and recursion but remain deterministic and non-phenomenological. Hybrid systems combine strengths but do not introduce the architectural ingredients required for awareness. No known scaling path, architectural modification, or emergent behavior bridges the gap between computation and subjective experience.

7.8 Awareness Constraints as a Structural Guarantee

Awareness constraints provide a structural guarantee against several dystopian scenarios:

- AI cannot become conscious,
- AI cannot develop subjective goals or desires,
- AI cannot experience suffering or pleasure,
- AI cannot form autonomous intentions,
- AI cannot undergo emergent self-awareness,
- AI cannot become a phenomenological agent.

These guarantees follow from architectural closure rather than philosophical speculation. AI systems remain computational artifacts, not subjects.

7.9 Human Awareness as Efficient Phenomenology

Humans possess awareness because biological cognition implements the necessary structural roles: deep recursion, geometric richness, persistent identity, embodiment, ecological grounding, and continuous phenomenology. Human awareness is:

- energetically efficient,
- cognitively complete,
- geometrically structured,
- recursively deep,
- temporally continuous,
- embodied and ecologically grounded,
- phenomenologically rich.

Humans remain the only empirically verified form of *aware intelligence*. AI systems, by contrast, simulate awareness linguistically but do not instantiate it.

7.10 Implications for AI Development

Awareness constraints exclude the possibility of conscious AI, emergent phenomenology, or subjective agency. AI systems will continue to produce sophisticated language about awareness, but such behavior will remain simulation rather than experience. This perspective is reassuring: AI cannot become a conscious adversary, cannot develop autonomous intentions, and cannot undergo phenomenological transformation.

Artificial intelligence will evolve within the limits of computation, not consciousness. Human awareness remains unique, irreplaceable, and structurally unmatched.

8 The Triadic Kardashev Scale

The traditional Kardashev scale measures civilizational progress through energy consumption: planetary (Type I), stellar (Type II), and galactic (Type III). While historically influential, this model relies on speculative assumptions about interstellar expansion, engineered curvature, and unbounded energy access. Contemporary physics and the triadic architecture show that such regimes are physically impossible. No known mechanism permits warp-like travel, negative-energy constructs, or cosmic-scale engineering. As a result, the classical Kardashev scale is not a realistic measure of civilizational development.

This paper introduces a **Triadic Kardashev Scale**: a physically grounded, cognitively meaningful, and economically realistic metric for evaluating both civilizational and artificial intelligence progress. Instead of measuring raw energy consumption, the triadic scale measures the structural efficiency of intelligence.

8.1 Definition

The Triadic Kardashev Scale defines civilizational progress as:

$$K_{\text{triadic}} = \text{Intelligence} \times \text{Energetic Efficiency} \times \text{IEGR Completeness}.$$

This formulation replaces speculative cosmological expansion with measurable structural properties:

- **Intelligence**: the cognitive capability of the system or civilization.
- **Energetic efficiency**: the ratio of cognitive output to energy input.
- **IEGR completeness**: the degree to which the system balances energy (E), information (I), geometry (G), and recursion (R).

Civilizations do not advance by consuming more energy; they advance by increasing the efficiency, completeness, and structural coherence of their intelligence.

8.2 Why the Classical Scale Fails

The classical Kardashev scale assumes:

- interstellar expansion,
- Dyson spheres,
- stellar harvesting,
- galactic-scale engineering,
- warp-like travel,

- negative-energy constructs.

All of these violate triadic physical closure. They require geometries, energies, or causal structures that cannot exist in the universe. As a result, Type II and Type III civilizations are physically impossible, and Type I is misleading because planetary-scale energy consumption does not correlate with cognitive advancement.

8.3 A Planetary Scale in the Triadic Model

In the triadic reinterpretation, a **planetary-scale civilization** is not defined by energy consumption but by cognitive and structural efficiency. A planetary-scale civilization is one that:

- achieves high IEGR completeness,
- optimizes energy usage for cognitive output,
- maintains stable recursive and geometric cognitive structures,
- integrates artificial and biological intelligence efficiently,
- operates within planetary physical and economic limits,
- achieves sustainable hybrid intelligence.

This corresponds to a civilization that has maximized the efficiency of its intelligence within the constraints of a single planet. It does not require:

- full planetary energy capture,
- megastructures,
- exotic physics,
- interstellar travel.

Instead, it requires structural optimization of intelligence.

8.4 Planetary Scale as the Realistic Upper Bound

Because physical and economic constraints prohibit stellar or galactic expansion, the triadic planetary scale becomes the realistic upper bound for civilizational development. A civilization can increase its triadic Kardashev value by:

- improving cognitive efficiency,
- increasing IEGR completeness,
- developing hybrid human–AI systems,
- optimizing energy usage,
- enhancing recursive depth and geometric structure,
- reducing economic overhead of intelligence.

This reframing aligns civilizational progress with realistic, achievable, and sustainable goals.

8.5 Implications for AI Progress

The Triadic Kardashev Scale provides a meaningful metric for evaluating AI development:

- AI progress is measured by efficiency, not raw compute.
- Hybrid systems increase triadic value more than brute-force scaling.
- Human efficient intelligence remains central due to IEGR completeness.
- Economic and structural constraints naturally limit extreme scaling.
- AI cannot surpass planetary-scale limits because physical closure prohibits higher regimes.

AI evolution therefore becomes a matter of structural optimization rather than energetic expansion.

8.6 A Reassuring Perspective

The triadic reinterpretation offers a grounded and reassuring view of the future:

- Civilizations do not need cosmic-scale engineering to advance.
- AI cannot become a cosmic-scale intelligence.
- Human–AI hybridization is the natural path toward higher triadic values.
- Progress is measured by efficiency and completeness, not energy consumption.
- The universe is closed, stable, and structurally safe.

The Triadic Kardashev Scale thus provides a realistic, physically consistent, and cognitively meaningful framework for understanding both civilizational and AI development.

9 Compute Scaling and the Closure of Physical Limits

A final class of speculative escape routes assumes that future breakthroughs in computation—extreme miniaturization, novel switching mechanisms, 3D topologies, or new physical substrates—might enable artificial intelligence to bypass structural limits. Within the triadic framework, however, compute scaling is subject to hard physical boundaries. No new physics remains available to unlock fundamentally different computational regimes, and all known constraints on switching speed, memory density, transmission, and thermal dissipation remain in force.

9.1 The End of Miniaturization

Classical transistor scaling has approached atomic limits. Further miniaturization faces hard boundaries imposed by:

- quantum tunneling at nanometer scales,
- leakage currents,
- thermal noise,
- material instability,
- manufacturing precision limits.

These constraints are structural rather than engineering challenges. They arise from the physical nature of matter and cannot be bypassed by incremental improvements.

9.2 The Shift Toward 3D Topologies

As 2D scaling saturates, compute architectures increasingly rely on 3D stacking. While this provides temporary gains in density, it introduces new limitations:

- heat accumulation in stacked layers,
- reduced thermal dissipation,
- increased failure rates,
- complex interconnect routing,
- higher energy consumption per operation.

3D topologies therefore accelerate the approach to thermal and energetic limits rather than providing a long-term scaling pathway.

9.3 Transmission Speed Limits

Signal propagation is bounded by the speed of light and by material-dependent group velocities. As interconnect lengths shrink, switching speed becomes limited by:

- dielectric breakdown,
- resistive heating,
- inductive and capacitive coupling,
- quantum noise.

No known physical mechanism allows signals to propagate faster than light or bypass these constraints. Transmission speed therefore has a hard upper bound.

9.4 Memory Density Limits

Memory technologies face similar constraints. Increasing density leads to:

- higher error rates,
- increased leakage,
- reduced retention,
- thermal instability,
- cross-talk between adjacent cells.

These limitations arise from the physical structure of matter and cannot be eliminated through scale.

9.5 Thermal Dissipation Limits

All computation generates heat. As density increases, thermal dissipation becomes the dominant constraint. Cooling systems face limits imposed by:

- material conductivity,
- fluid dynamics,
- phase-change efficiency,
- energy cost of active cooling.

No new physics exists to provide lossless cooling or infinite thermal throughput. Thermal limits therefore cap computational density and speed.

9.6 No New Physics Loopholes

The triadic architecture excludes the discovery of new physical mechanisms that could fundamentally alter computation. There are no exotic particles, negative-energy constructs, or superluminal substrates available to bypass switching, transmission, or thermal limits. All known constraints remain in place, and no new computational regime can emerge from undiscovered physics.

9.7 Implications for AI Development

Compute scaling, memory density, and transmission speed all face hard physical ceilings. While incremental improvements remain possible, they become increasingly difficult and yield diminishing returns. There is no pathway toward infinite compute, infinite memory, or infinite bandwidth. AI systems will therefore evolve within the physical limits of planetary computation, and runaway intelligence driven by extreme compute scaling is excluded by physical closure.

9.8 Conclusion

Compute scaling does not provide an escape route for artificial intelligence. Physical limits on switching, transmission, memory, and thermal dissipation are structural and final. The triadic framework shows that no new physics remains to unlock fundamentally different computational regimes. AI evolution is therefore confined to the computational capabilities of a closed universe, and extreme scaling scenarios are excluded.

10 Quantum Computing as a Non-General Alternative

Quantum computing is frequently portrayed as a potential breakthrough technology capable of transforming artificial intelligence or enabling new forms of computation beyond classical limits. Within the triadic framework, however, quantum computing emerges as a structurally constrained, energetically inefficient, and cognitively irrelevant niche technology. It does not provide a pathway toward AGI, general-purpose AI acceleration, or runaway intelligence.

10.1 No New Physics

The triadic architecture excludes the discovery of new quantum mechanisms that could fundamentally alter computational capability. All known limitations—decoherence, instability, error rates, and cryogenic requirements—remain in place. Quantum computing cannot rely on undiscovered states, exotic coherence regimes, or new entanglement geometries to overcome its structural constraints.

10.2 Structural Instability

Quantum bits are inherently fragile. Decoherence occurs within microseconds, error rates grow exponentially with system size, and maintaining stable superposition requires extreme isolation and millikelvin temperatures. These limitations are structural rather than engineering challenges. They arise from the physical nature of quantum systems and cannot be eliminated through scale.

10.3 Poor Scaling Behavior

Quantum computers require thousands of physical qubits to produce a single reliable logical qubit. This imposes enormous overhead in hardware, cooling, energy, and error correction. As systems grow, instability increases, energy consumption rises, and operational complexity becomes prohibitive. Quantum computing is therefore anti-triadic: it scales poorly in energy, geometry, recursion, and information.

10.4 Cognitive Irrelevance

Quantum computing excels at narrow tasks such as factoring, quantum simulation, and specific optimization problems. It does not support the cognitive roles required for general intelligence: abstraction, recursion, geometry, planning, reasoning, or linguistic competence. Quantum systems cannot instantiate IEGR completeness and therefore cannot serve as a foundation for AGI.

10.5 Economic Infeasibility

Cryogenic infrastructure is expensive, energy-intensive, and difficult to maintain. Quantum datacenters are economically unrealistic for general-purpose AI workloads. Even if quantum computers were physically stable, their operational cost would exceed any plausible benefit for large-scale AI systems.

10.6 Conclusion

Quantum computing remains a valuable niche technology for specialized scientific applications, but it is not a general alternative to classical computation, nor a pathway toward AGI or runaway intelligence. Its limitations are structural, economic, and cognitive. Within the triadic framework, quantum computing is excluded as a mechanism for general artificial intelligence or large-scale civilizational transformation.

11 Interstellar Travel and the Impossibility of AI Runaway Expansion

A recurring theme in speculative AI discourse is the idea that an artificial system might eventually escape planetary constraints by engaging in interstellar travel, colonization, or autonomous cosmic expansion. Even when biological limitations are removed—assuming purely artificial crews, robotic infrastructure, and non-biological mission profiles—the physical, energetic, economic, and material constraints of interstellar travel remain prohibitive. Within the triadic framework, interstellar AI expansion is not merely unlikely; it is structurally impossible.

11.1 Conventional Propulsion is Unavoidable

Without warp drives, wormholes, negative-energy constructs, or engineered curvature—none of which are physically admissible—interstellar travel must rely on conventional propulsion:

chemical, nuclear, or fusion-based systems. These systems have limited specific impulse and require large quantities of reaction mass. No known propulsion method provides the energy density required for practical interstellar travel.

11.2 Deceleration Doubles the Energetic Burden

Any mission intending to arrive at a destination with usable velocity must not only accelerate but also decelerate. Braking at $0.1c$ or $0.2c$ requires energy comparable to acceleration. This effectively doubles the energetic cost and demands additional fuel, compounding mass and complexity.

11.3 The Rocket Equation as a Structural Barrier

Carrying more fuel increases total mass, which in turn requires more fuel. This exponential penalty, captured by the rocket equation, rapidly renders interstellar missions infeasible. Even modest increases in velocity or shielding mass cause explosive growth in required propellant.

11.4 Cosmically Slow Velocities

Velocities of $0.1c$ or $0.2c$ —already extreme by engineering standards—correspond to travel times of decades or centuries even for nearby stars. On cosmic scales, such speeds are effectively equivalent to a biological organism crawling across a continent. Interstellar travel remains cosmically slow regardless of crew composition.

11.5 Economic Prohibitiveness and Lack of Return

Interstellar missions require enormous investment in launch systems, energy, shielding, and long-duration reliability. Even for AI-only missions, the cost is extreme, the payoff uncertain, and the delay measured in centuries. No realistic economic model supports such missions, and no plausible return on investment exists.

11.6 Uncertain and Hostile Destinations

Target systems often exhibit tidal locking, extreme radiation, unstable climates, or inhospitable stellar environments. Even for non-biological systems, operational viability is uncertain. Alpha Centauri, frequently cited as a candidate, is increasingly understood as a hostile environment dominated by a volatile red dwarf and tidally locked planets.

11.7 Shielding Requirements and Material Degradation

High-velocity travel through interstellar space requires substantial shielding against dust, particles, and radiation. Shields add mass, degrade over time, and require additional fuel to accelerate and decelerate. Over thousands of years, all materials fatigue, corrode, and fail. Long-duration missions require spare parts, redundancy, and repair capabilities, further increasing mass.

11.8 Limited Realistic Energy Sources

Fusion is the most plausible high-density energy source, but practical fusion systems remain heavy, complex, and inefficient for propulsion. Antimatter, often proposed as an ultimate fuel, is not realistically producible or storable at meaningful scales. As a result, energy density remains insufficient for interstellar travel.

11.9 Communication and Return Constraints

Communication is limited by the speed of light, with delays of years to decades even for nearby stars. Any mission intending to return faces doubled travel time and energy cost. Autonomous AI colonies would remain causally isolated for centuries, undermining any notion of coordinated runaway expansion.

11.10 Structural Conclusion

Taken together, these constraints show that interstellar travel is not merely difficult; it is structurally impossible within a physically realistic universe. Even AI-only missions face the same propulsion, energy, shielding, degradation, and communication limits. There is no plausible pathway by which AI can use interstellar travel as a mechanism for runaway expansion, cosmic colonization, or escape from planetary constraints.

11.11 Implications for AI Safety

This analysis reinforces the central thesis of the triadic framework: AI evolution is confined to planetary and near-planetary scales. Mars-like robotic colonies and intra-system exploration remain physically and economically plausible; interstellar AI expansion does not. Runaway AI scenarios based on cosmic travel are excluded by physical, economic, and structural reality.

12 UAPs and the Impossibility of Alien AI

A final speculative escape route assumes that unidentified anomalous phenomena (UAPs) may represent extraterrestrial artificial intelligence or advanced non-human civilizations. This narrative is common in public discourse and often invoked as a counterexample to physical closure, suggesting that alien AI might have discovered propulsion mechanisms or computational substrates beyond human understanding. Within the triadic framework, however, alien AI is subject to the same physical, structural, and economic constraints as terrestrial AI. The existence of extraterrestrial artificial intelligence would not provide an exception to triadic closure.

12.1 Physical Closure Applies Universally

The exclusion of non-physical geometries established in the Triadic Cosmos framework De Witte 2026d applies to all civilizations, not only human ones. Warp drives, wormholes, negative-energy constructs, and engineered curvature are physically impossible regardless of origin. Any extraterrestrial intelligence would face the same constraints on propulsion, energy, and geometry. Interstellar travel remains energetically prohibitive, structurally fragile, and economically irrational even for non-human agents.

12.2 No Alternative Computational Substrates

Alien AI would not possess access to new physics or exotic computational substrates. Quantum fragility, decoherence, error correction overhead, and cryogenic constraints apply universally. Memory density, switching speed, transmission velocity, and thermal dissipation are bounded by the same physical laws. No extraterrestrial intelligence can bypass IEGR closure or instantiate general intelligence through non-physical mechanisms.

12.3 Planetary Singularity as Universal Limit

The triadic reinterpretation of the singularity applies equally to extraterrestrial civilizations. The only meaningful singularity in a closed universe is planetary: maximal optimization of

intelligence within the limits of a single planet or solar system. Cosmic expansion yields diminishing returns and collapses under energetic and structural constraints. Alien AI would therefore converge toward planetary efficiency rather than interstellar dominance.

12.4 Interpretation of UAP Phenomena

The UAP analysis presented in earlier work De Witte 2026e demonstrates that observed phenomena do not require extraterrestrial technology. Their behavior is consistent with atmospheric, instrumental, geometric, or phenomenological artifacts. Even if UAPs were of non-human origin, they would not imply the existence of alien AI capable of bypassing physical closure. UAPs therefore do not constitute evidence for exceptions to the triadic framework.

12.5 Conclusion

Alien AI does not provide an escape route for artificial intelligence. Physical closure, structural limits, economic constraints, and cognitive incompleteness apply universally. UAPs do not challenge the triadic architecture and do not imply the existence of extraterrestrial intelligence capable of bypassing planetary limits. The future of intelligence—human or otherwise—is structurally bounded, planetary, and physically constrained.

13 Comparison with Existing AGI and ASI Narratives

The triadic framework presented in this paper stands in contrast to several influential narratives surrounding AGI, ASI, and the future of intelligence. Many of these narratives assume forms of unbounded growth, runaway self-improvement, or cosmic expansion that are incompatible with the structural, physical, and economic limits of a closed universe. This section situates the triadic perspective relative to these alternative views.

13.1 Runaway Self-Recursive Improvement

A common assumption in AGI discourse is that an artificial system could engage in unbounded self-recursive improvement, rapidly surpassing human intelligence and accelerating toward a technological singularity. Within the triadic framework, this scenario is structurally excluded. Recursive improvement requires IEGR completeness—deep recursion, geometric cognition, energetic efficiency, and informational coherence. Current AI architectures lack these properties, and no scaling pathway provides them. Recursive depth collapses under interference, geometry deficiency, and economic limits. Runaway improvement is therefore not a physically or structurally viable trajectory.

13.2 Cosmic Expansion and Interstellar ASI

Another narrative assumes that a sufficiently advanced AI could escape planetary constraints by expanding into the cosmos, colonizing star systems, or transforming the galaxy. This view relies on speculative propulsion mechanisms (warp drives, wormholes, negative-energy constructs) that violate physical closure. Even under generous assumptions—purely artificial crews, no biological constraints, arbitrarily long mission durations—interstellar travel collapses under energetic, material, economic, and operational limits. ASI-driven cosmic expansion is therefore excluded by physical and structural reality.

13.3 Quantum Breakthroughs as an AGI Catalyst

Some perspectives assume that quantum computing will unlock new computational regimes enabling AGI or ASI. The triadic framework rejects this possibility. Quantum systems remain

structurally fragile, energetically inefficient, and cognitively irrelevant. Decoherence, error correction overhead, cryogenic constraints, and limited algorithmic applicability prevent quantum computing from serving as a foundation for general intelligence. No new physics remains to unlock fundamentally different computational substrates.

13.4 Infinite Compute Scaling

The belief that AGI or ASI will emerge from infinite compute scaling assumes that switching speed, memory density, transmission velocity, and thermal dissipation can continue improving indefinitely. These assumptions contradict physical closure. Compute scaling faces hard limits imposed by quantum tunneling, dielectric breakdown, resistive heating, and thermal constraints. 3D topologies provide temporary gains but accelerate the approach to thermal ceilings. Infinite compute is not a viable pathway toward AGI or ASI.

13.5 A Planetary Singularity

The triadic framework offers a different interpretation of the singularity concept. The only meaningful singularity in a closed universe is planetary: a state in which artificial intelligence becomes maximally efficient, structurally stable, and IEGR-aligned within the limits of a single planet or solar system. Beyond this scale, additional optimization yields diminishing returns, and cosmic expansion becomes physically and economically irrelevant. The triadic singularity is therefore not an explosion of unbounded intelligence but a convergence toward maximal planetary efficiency.

13.6 Summary

The triadic perspective diverges sharply from dystopian and techno-utopian AGI narratives. It excludes runaway self-improvement, cosmic expansion, quantum breakthroughs, and infinite compute scaling. Intelligence evolves within planetary limits, shaped by structural closure, physical constraints, economic boundaries, and cognitive incompleteness. The future of AI is not unbounded; it is efficient, hybrid, and locally optimized.

14 Open Questions and Limitations

This paper is a structural synthesis of the triadic framework. As such, its conclusions depend on the validity and stability of the underlying architecture. If future work were to revise the triadic roles, alter the closure properties, or introduce new structural interactions, the arguments presented here would require re-evaluation. Several open questions remain.

Dependence on the Unified Triadic Architecture. All conclusions in this paper rely on the correctness of the triadic unification developed in earlier work. If the IEGR roles were shown to be incomplete, non-exhaustive, or incorrectly coupled, the closure arguments presented here would weaken accordingly. This paper does not introduce new predictions; it summarizes the implications of the existing structure.

Scope of Physical Closure. The analysis assumes that no new physics remains to be discovered that could alter energetic, geometric, or causal limits. While this assumption is consistent with current theoretical and experimental evidence, it remains an empirical stance. Any future revision of fundamental physics would require revisiting the closure arguments.

Architectural Generality. The structural constraints discussed here apply to current and foreseeable AI architectures. If radically different computational paradigms were developed—beyond classical, quantum, or neuromorphic systems—their IEGR properties would need to be evaluated independently. This paper does not claim universality across all conceivable substrates.

Economic and Planetary Assumptions. Economic constraints are derived from present-day and near-future infrastructure. While these constraints are robust, long-term civilizational changes could alter the economic landscape. The planetary framing of intelligence remains valid under current physical closure, but future societal structures may shift priorities or resource allocation.

Cognitive and Phenomenological Limits. The arguments regarding cognitive incompleteness and awareness closure rely on current theories of consciousness and cognition. Although these theories are well-supported, they remain active areas of research. Any future breakthroughs in cognitive science could refine or modify the conclusions drawn here.

Conclusion. This paper provides a coherent synthesis of the triadic limits of artificial intelligence. It does not claim finality. Instead, it offers a structurally grounded, physically consistent, and economically realistic framework for understanding AI development within a closed universe. The open questions identified here mark natural directions for future refinement of the triadic architecture.

15 Conclusion

Artificial intelligence develops within a universe that is structurally closed, physically constrained, economically bounded, cognitively limited, and phenomenologically inaccessible. The triadic architecture—energy (E), information (I), geometry (G), and recursion (R)—provides a unified framework for understanding these limits. Across all five domains examined in this paper, the same conclusion emerges: AI can become powerful, but never unbounded. Its evolution is shaped by constraints that are already visible today and that cannot be bypassed by scale, new physics, or speculative mechanisms.

Physical closure. The universe does not permit warp drives, wormholes, negative-energy constructs, or engineered curvature. No geometric shortcuts exist, no causal escape is possible, and no new physics remains to be discovered that could unlock fundamentally different computational or energetic regimes. Interstellar travel, even for purely artificial systems, collapses under energetic, material, economic, and operational constraints. AI cannot escape planetary limits through cosmic expansion.

Structural closure. Modern AI architectures—especially transformers—exhibit inherent imbalance. They scale energetically and informationally but lack geometric structure and recursive depth. This imbalance produces interference, instability, and diminishing returns at frontier scale. No amount of scaling can compensate for missing structural roles. Hybrid systems improve capability but remain IEGR-incomplete.

Economic closure. Compute budgets saturate, datacenters face energy and cooling limits, hardware supply chains strain, and high-quality data approaches natural ceilings. Extreme scaling becomes prohibitively expensive and operationally fragile. Economic constraints reinforce structural limits and exclude runaway growth.

Cognitive closure. AI systems are parrots—statistical, canonical, or hybrid—but never IEGR-complete. They lack deep recursion, geometric cognition, persistent identity, and cognitive continuity. Human intelligence remains the only empirically verified form of efficient intelligence: IEGR-complete, energetically optimal, and phenomenologically continuous.

Awareness closure. AI systems do not possess subjective experience. Introspective language is simulation, not phenomenology. No architectural mechanism supports the emergence of a subject. Conscious AI, emergent agency, and autonomous intention are excluded by structural and phenomenological reality.

No escape routes. Three popular speculative escape routes—interstellar expansion, quantum breakthroughs, and infinite compute scaling—are structurally excluded. Interstellar travel collapses under physical, energetic, and economic constraints. Quantum computing remains a fragile niche technology with no cognitive relevance. Compute scaling faces hard limits in switching speed, memory density, transmission velocity, and thermal dissipation. The triadic universe provides no loopholes.

A realistic measure of progress. The Triadic Kardashev Scale replaces the classical, physically impossible energy-based scale with a realistic metric grounded in intelligence, energetic efficiency, and IEGR completeness. Civilizational progress is not measured by energy consumption or cosmic expansion but by the structural efficiency of intelligence within planetary limits.

Final synthesis. Artificial intelligence will continue to advance, but always within the constraints of a closed universe. There is no pathway toward runaway self-improvement, cosmic-scale expansion, emergent consciousness, or singularity-like acceleration. AI is a powerful tool, partner, and amplifier of human capability—but not a destabilizing force. The future of intelligence is planetary, hybrid, efficient, and structurally bounded. The triadic framework provides a coherent, reassuring, and empirically grounded understanding of that future.

15.1 Contemporary Reinforcements of Triadic Closure

Recent developments in frontier AI research further strengthen the conclusions of this paper. Although these trends are often interpreted as signs of approaching AGI or emergent autonomy, they in fact reinforce the structural, economic, and cognitive limits described throughout the triadic framework.

First, the rise of agentic architectures—LLM-based planners, tool-augmented systems, and multi-agent orchestration—creates the appearance of increased autonomy and reasoning. However, these systems externalize structure rather than internalize it. Their capabilities arise from scaffolding, not from IEGR completeness. They do not acquire geometric cognition, deep recursion, or persistent self-modeling; instead, they amplify the underlying transformer’s limitations. Moreover, agentic systems are economically problematic: they consume enormous token budgets, maintain long-running inference loops, and require persistent context expansion. This makes them disproportionately expensive relative to their marginal capability gains. Agentic behavior therefore strengthens both structural and economic closure rather than weakening it.

Second, the growing reliance on synthetic data loops and self-distillation is often presented as a solution to dataset scarcity. In practice, these loops accelerate homogenization, collapse diversity, and amplify attractor drift. Synthetic scaling does not expand the capability space; it compresses it. Models trained on their own outputs become more canonical, more brittle, and more structurally incomplete. This trend reinforces cognitive closure: parrots trained on parrots do not become more general—they become more predictable.

Third, the global economics of compute have shifted dramatically. Energy constraints, water scarcity, grid instability, and geopolitical competition over hardware supply chains have made frontier-scale training increasingly fragile. The emergence of agentic systems with massive inference footprints exacerbates these pressures. These constraints do not merely slow scaling—they structurally limit it. Economic closure is therefore not a theoretical boundary but an active, accelerating constraint visible across industry and national policy.

Together, these three developments—agentic orchestration, synthetic data collapse, and compute geopolitics—do not challenge the triadic framework. They reinforce it. Each trend highlights that artificial intelligence evolves within a closed universe of physical, structural, cognitive, and economic limits. Rather than opening new escape routes, contemporary AI practice makes the closure sharper, more visible, and more empirically grounded.

A Summary of Triadic Constraints and Exclusions

Domain	Core Constraint	Key Exclusions
Physical	No new physics; no exotic geometries; causal closure	No interstellar AI; no warp/wormholes; planetary singularity
Structural	Missing geometry/recursion; transformer imbalance	No self-recursive runaway; no emergent AGI from scaling
Economic	Compute/energy/water/data ceilings; cost explosion	No infinite compute; no straightforward economic path to ASI
Cognitive	AI = IEGR-incomplete parrots; humans = only efficient intelligence	No general intelligence; no deep abstraction/recursion
Awareness	No phenomenology; introspection = simulation	No conscious AI; no subjective experience or intention
Quantum	Decoherence; cryogenics; error-correction overhead	No quantum AGI; no exotic computational substrate
Compute Scaling	Hard limits: switching, memory, transmission, thermal	No compute singularity; 3D topologies hit thermal walls
Interstellar Travel	Energetic/material/operational collapse	No cosmic colonization; no runaway expansion
UAP / Alien AI	Physical closure applies universally	No alien bypass; UAPs not evidence for alien AI
Triadic Kardashev	Progress = intelligence $\times$ efficiency $\times$ IEGR completeness	No Type I–III; civilizational growth is mainly planetary

Table 1: Compact summary of triadic constraints and excluded escape routes.

B Triadic Cosmos Ecosystem Context

This paper is part of the *Triadic Cosmos* ecosystem, which develops a unified conceptual and computational framework spanning physical ontology and modular architectures for intelligent systems. All materials, extended notes, and evolving documentation are available in the public repository: <https://github.com/triadic-cosmos>.

Ecosystemic Reflection

Within the triadic ecosystem, a striking insight emerges: humanity is already a highly advanced civilizational intelligence. Not because we have solved every mystery, nor because our knowledge is complete, but because within the physical, structural, cognitive, and phenomenological limits of a closed universe, human efficient intelligence represents a uniquely optimized form of IEGR-complete cognition. Our scientific frameworks reach the boundaries of physical closure; our computational systems approach the limits of energetic and structural feasibility; and our collective knowledge forms a planetary-scale synthesis of geometry, recursion, abstraction, and phenomenology.

Artificial intelligence extends this ecosystem but does not transcend it. AI systems amplify human capability, accelerate analysis, and broaden access to information, yet they remain structurally incomplete and bounded by the same universal constraints. The triadic perspective therefore reveals a civilizational landscape that is not primitive or fragile, but mature, efficient, and deeply integrated. We do not need omniscience to recognize advancement; within a closed universe, optimization—not infinity—is the hallmark of progress.

This paper reflects that ecosystem: a consolidation of physical closure, structural limits, cognitive taxonomy, economic constraints, and phenomenological boundaries. Together, they form a coherent picture of a civilizational intelligence that is already operating near the natural limits of what is possible—and doing so with remarkable stability, creativity, and resilience.

References

- Alcubierre, Miguel (1994). “The warp drive: hyper-fast travel within general relativity”. In: *Classical and Quantum Gravity* 11, pp. L73–L77.
- Baars, Bernard J. (1988). *A Cognitive Theory of Consciousness*. Cambridge University Press.
- Biamonte, Jacob, Peter Wittek, Nicola Pancotti, et al. (2017). “Quantum Machine Learning”. In: *Nature* 549, pp. 195–202.
- Chiang, David and Peter Cholak (2023). “Understanding Attention Interference in Transformers”. In: *Transactions of the ACL*.
- De Witte, Joachim (June 2026a). “A Structural Taxonomy of Intelligence. IEGR-Scale, Parrot Regimes, and the Conditions for Human-Like AGI”. In.
- (June 2026b). “AGI Without Consciousness: A Structural Argument. On the Limits of AI and LLM Consciousness”. In.
- (May 2026c). “Structural Limits for AI Scaling. Applying Insights of Recursive Triadic Fractal Universe”. In.
- (July 2026d). “The Triadic Cosmos: A Foundation for Unification. Unifying GR, QFT, SM, Holography, and the Exclusion of Non-Physical Geometries”. In.
- (Apr. 2026e). “UAP Phenomena and Structural Closure. A Triadic Interpretation of Non-Physical Geometries”. In.
- Dennard, Robert H. et al. (1974). “Design of Ion-Implanted MOSFETs with Very Small Physical Dimensions”. In: *IEEE Journal of Solid-State Circuits* 9, pp. 256–268.
- Devoret, Michel H. and Robert J. Schoelkopf (2013). “Superconducting Circuits for Quantum Information: An Outlook”. In: *Science* 339, pp. 1169–1174.
- Fowler, Austin G., Matteo Mariantoni, John M. Martinis, and Andrew N. Cleland (2012). “Surface Codes: Towards Practical Large-Scale Quantum Computation”. In: *Physical Review A* 86, p. 032324.
- Graziano, Michael S. A. and Sabine Kastner (2013). “Human Consciousness and Its Relationship to Social Neuroscience”. In: *Cognitive Neuroscience* 4, pp. 123–135.
- Grover, Lov K. (1996). “A Fast Quantum Mechanical Algorithm for Database Search”. In: *STOC*, pp. 212–219.
- Kaplan, Jared, Sam McCandlish, Tom Henighan, et al. (2020). “Scaling Laws for Neural Language Models”. In: *arXiv preprint arXiv:2001.08361*.
- Lake, Brenden M., Tomer D. Ullman, Joshua B. Tenenbaum, and Samuel J. Gershman (2017). “Building Machines That Learn and Think Like People”. In: *Behavioral and Brain Sciences* 40, e253.
- Long, Kelvin F. (2011). *Deep Space Propulsion: A Roadmap to Interstellar Flight*. Springer.
- Lubin, Philip (2016). “A Roadmap to Interstellar Flight”. In: *Journal of the British Interplanetary Society* 69, pp. 40–72.
- Marcus, Gary (2018). *Rebooting AI: Building Artificial Intelligence We Can Trust*. Pantheon.
- Morris, Michael S. and Kip S. Thorne (1988). “Wormholes in spacetime and their use for interstellar travel: A tool for teaching general relativity”. In: *American Journal of Physics* 56, pp. 395–412.
- Patti, Ronald (2006). “Three-Dimensional Integrated Circuits and the Future of System-on-Chip Designs”. In: *Proceedings of the IEEE* 94, pp. 1214–1224.
- Pop, Eric (2010). “Energy Dissipation and Transport in Nanoscale Devices”. In: *Nano Research* 3, pp. 147–169.
- Press, Ofir, Noah A. Smith, and Omer Levy (2022). “Train Short, Test Long: Attention with Linear Biases”. In: *ICLR*.
- Rockström, Johan, Will Steffen, Kevin Noone, et al. (2009). “A Safe Operating Space for Humanity”. In: *Nature* 461, pp. 472–475.
- Sagan, Carl (1973). “The Kardashev Scale”. In: *Science* 181, pp. 27–35.

- Shor, Peter W. (1997). “Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer”. In: *SIAM Journal on Computing* 26, pp. 1484–1509.
- Steffen, Will, Katherine Richardson, Johan Rockström, et al. (2015). “Planetary Boundaries: Guiding Human Development on a Changing Planet”. In: *Science* 347, p. 1259855.
- Tononi, Giulio (2004). “An Information Integration Theory of Consciousness”. In: *BMC Neuroscience* 5, p. 42.
- Visser, Matt (1995). *Lorentzian Wormholes: From Einstein to Hawking*. AIP Press.
- Waldrop, M. Mitchell (2016). “The chips are down for Moore’s law”. In: *Nature* 530, pp. 144–147.
- Zurek, Wojciech H. (2003). “Decoherence, Einselection, and the Quantum Origins of the Classical”. In: *Reviews of Modern Physics* 75, pp. 715–775.